

# Efficient Black-box Checking of Snapshot Isolation in Databases

(Conference VLDB'2024)

Hengfeng Wei

hfwei@nju.edu.cn

August 7, 2023



# Database Transactions

A database transaction is a *group* of operations



that should be executed **atomically**.

# Isolation Levels

Transactions may be executed concurrently.


The isolation levels specify how they are isolated from each other.

# Serializability (SER)

All transactions appear to execute serially, one after another.

too expensive, especially for distributed transactions

# Snapshot Isolation (SI)



# Snapshot Isolation (SI)

example

**Snapshot Read:** Each transaction reads data from a snapshot of committed data valid as of the (logical) time the transaction started.

**Snapshot Write:** Concurrent transactions cannot write to the same key. One of them must be aborted.

# SI Prevents the “Lost Update” Anomaly

$$T_0$$
$$\boxed{W(acct, 0)}$$

# SI Prevents the “Lost Update” Anomaly

$T_A$

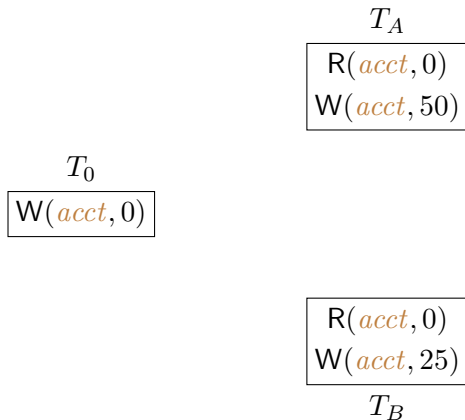
$R(acct, 0)$   
 $W(acct, 50)$

$T_0$

$W(acct, 0)$




# SI Prevents the “Lost Update” Anomaly



$T_A$  and  $T_B$  are executed concurrently.

# SI Prevents the “Lost Update” Anomaly



$T_A$  and  $T_B$  are executed concurrently.

# SI Prevents the “Causality Violation” Anomaly

$$T_A \boxed{W(x, post)}$$

# SI Prevents the “Causality Violation” Anomaly

$$T_A \boxed{W(x, post)}$$

$$T_B \boxed{\begin{array}{l} R(x, post) \\ W(y, comment) \end{array}}$$

# SI Prevents the “Causality Violation” Anomaly

$$T_A \boxed{W(x, post)}$$

$$T_B \boxed{\begin{array}{l} R(x, post) \\ W(y, comment) \end{array}}$$

$$T_C \boxed{\begin{array}{l} R(x, empty) \\ R(y, comment) \end{array}}$$

# SI Allows the “Write Skew” Anomaly

# Databases that Claim to Support SI

database logos

# Snapshot Isolation (SI)

Database systems may fail to provide SI as they claim.  
+papers



# The SI Checking Problem

Given a history  $H$  of a database system,  
to decide whether  $H$  satisfies SI?  
+fig

# Motivation: Black-box SI Checker

Since the internals of database systems are often unavailable or are hard to understand,  
a *black-box* SI checker is highly desirable.

# Motivation: Black-box SI Checker

A black-box SI checker should be

**Sound:** return no false positives

**Complete:** miss no violations

**Efficient:** run in a reasonable time even for large workloads

**Informative:** report understandable counterexamples

# Motivation: Black-box SI Checker

related-work

# Contributions: PolySI

Sound & Complete: characterization of SI in terms of *generalized polygraphs*

Efficient: encoding into MonoSAT queries;  
domain-specific pruning before solving


Informative: extract understandable counterexamples from  
MonoSAT unsatisfiable core

# Contributions: PolySI

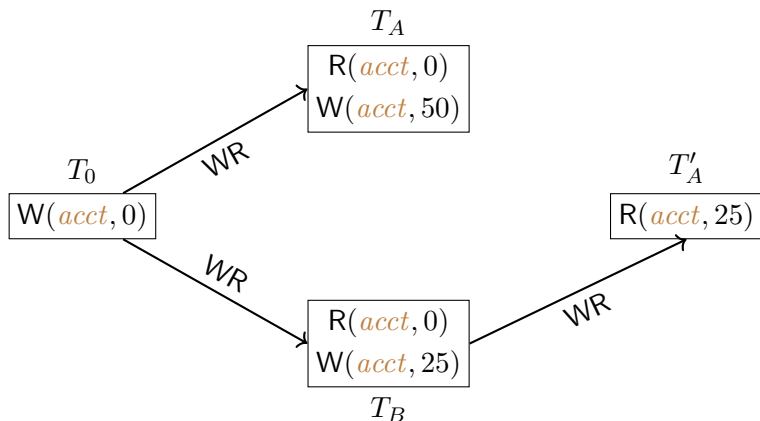
PolySI found SI violations in production database systems.

PolySI outperformed state-of-the-art black-box SI checkers and scales up to large workloads.

# Characterization of Snapshot Isolation



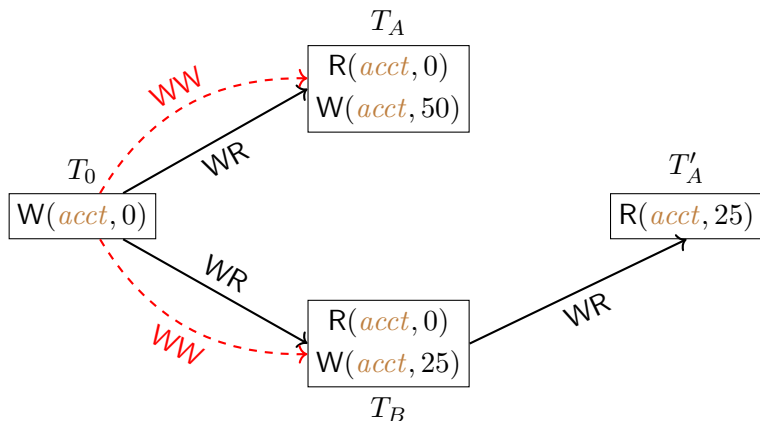
# Characterization of Snapshot Isolation



WR: “write-read” dependency capturing the “read-from” relation



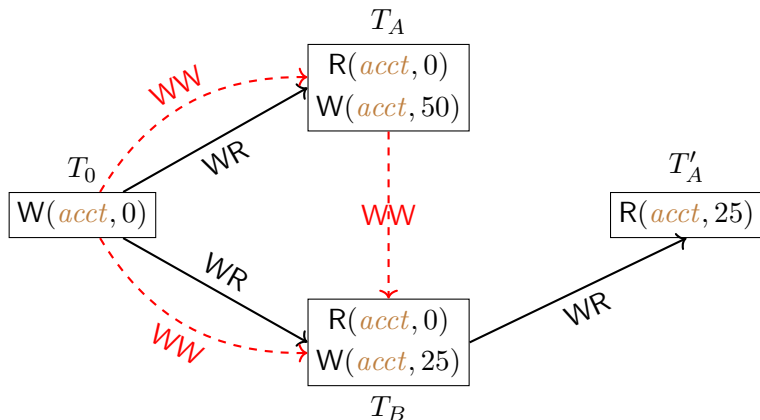
# Characterization of Snapshot Isolation



WW: “write-write” dependency capturing the version order

# Characterization of Snapshot Isolation

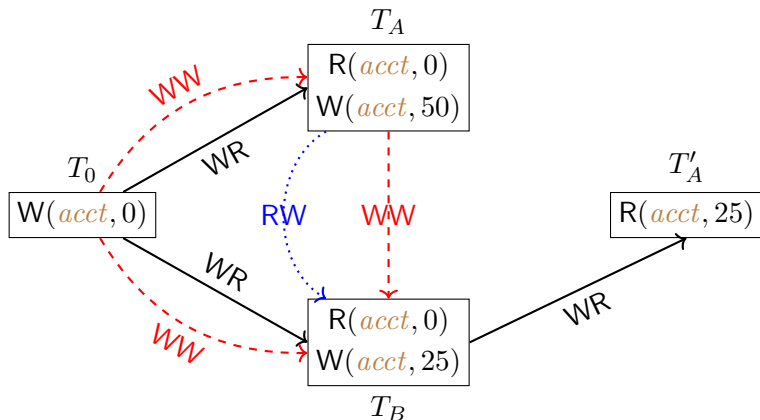
Suppose that  $T_A \xrightarrow{WW} T_B$



WW: “write-write” dependency capturing the version order

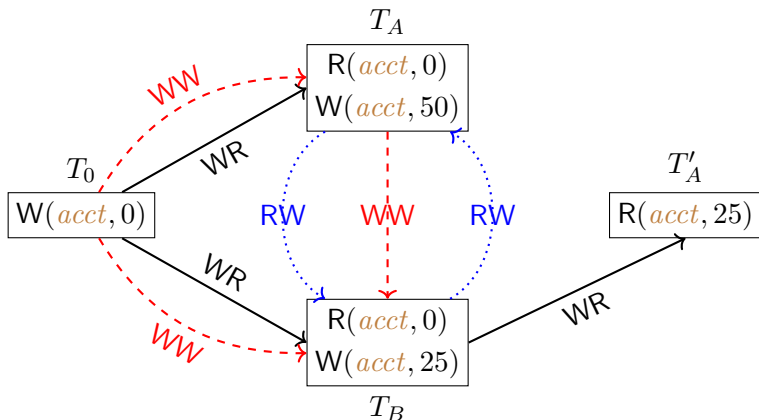
# Characterization of Snapshot Isolation

Suppose that  $T_A \xrightarrow{WW} T_B$



RW: “read-write” dependency capturing the overwritten relation

# Characterization of Snapshot Isolation



undesiable cycle:  $T_A \xrightarrow{WW} T_B \xrightarrow{RW} T_A$

# Characterization of Snapshot Isolation





Hengfeng Wei (hfwei@nju.edu.cn)