

# Efficient Black-box Checking of Snapshot Isolation in Databases

(Conference VLDB'2024)

Hengfeng Wei

hfwei@nju.edu.cn

August 9, 2023



# Database Transactions

A database transaction is a *group* of operations



that should be executed **atomically**.

# Isolation Levels

Transactions may be executed concurrently.


The isolation levels specify how they are isolated from each other.

# Serializability (SER)

All transactions appear to execute serially, one after another.

too expensive, especially for distributed transactions

# Snapshot Isolation (SI)



# Snapshot Isolation (SI)

example

**Snapshot Read:** Each transaction reads data from a snapshot of committed data valid as of the (logical) time the transaction started.

**Snapshot Write:** Concurrent transactions cannot write to the same key. One of them must be aborted.

# SI Prevents the “Lost Update” Anomaly

$$T_0$$
$$\boxed{W(acct, 0)}$$

# SI Prevents the “Lost Update” Anomaly

$T_A$

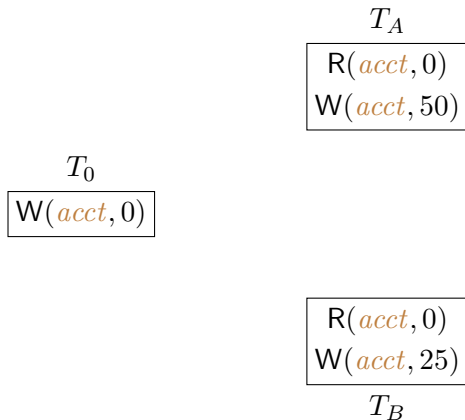
$R(acct, 0)$   
 $W(acct, 50)$

$T_0$

$W(acct, 0)$




# SI Prevents the “Lost Update” Anomaly



$T_A$  and  $T_B$  are executed concurrently.

# SI Prevents the “Lost Update” Anomaly



$T_A$  and  $T_B$  are executed concurrently.

# SI Prevents the “Causality Violation” Anomaly

$$T_A \boxed{W(x, post)}$$

# SI Prevents the “Causality Violation” Anomaly

$$T_A \boxed{W(x, post)}$$

$$T_B \boxed{\begin{array}{l} R(x, post) \\ W(y, comment) \end{array}}$$

# SI Prevents the “Causality Violation” Anomaly

$$T_A \boxed{W(x, post)}$$

$$T_B \boxed{\begin{array}{l} R(x, post) \\ W(y, comment) \end{array}}$$

$$T_C \boxed{\begin{array}{l} R(x, empty) \\ R(y, comment) \end{array}}$$

# SI Allows the “Write Skew” Anomaly

# Databases that Claim to Support SI

database logos

# Snapshot Isolation (SI)


Database systems may fail to provide SI as they claim.  
+papers



# The SI Checking Problem

## Definition (The SI Checking Problem)


The SI checking problem is the **decision problem** of determining whether a given **history**  $\mathcal{H} = (T, SO)$  satisfies SI?




# Motivation: Black-box SI Checker

Since the internals of database systems are often unavailable or are hard to understand,  
a *black-box* SI checker is highly desirable.

# Motivation: Black-box SI Checker




# Motivation: Black-box SI Checker




*Sound:* If the checker says **X**, then the history is not SI.

# Motivation: Black-box SI Checker




*Complete:* If the checker says ✓, then the history is SI.

# Motivation: Black-box SI Checker



*Efficient:* The checker should scale up to large workloads.

# Motivation: Black-box SI Checker




*Informative:* The checker should provide understandable counterexamples if it says **X**.

# Motivation: Black-box SI Checker

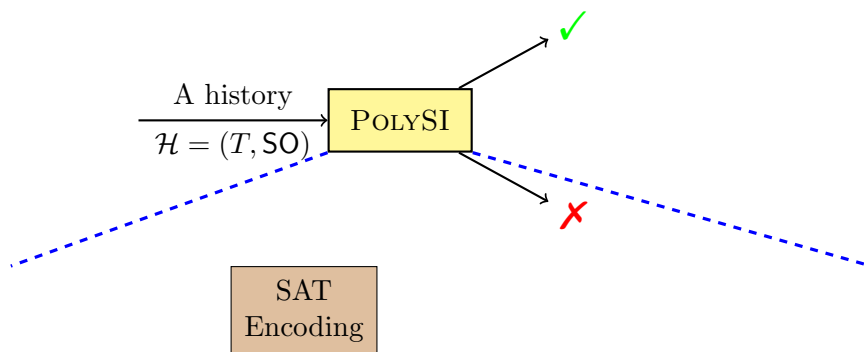
related-work



# Contributions: the POLYSI Checker

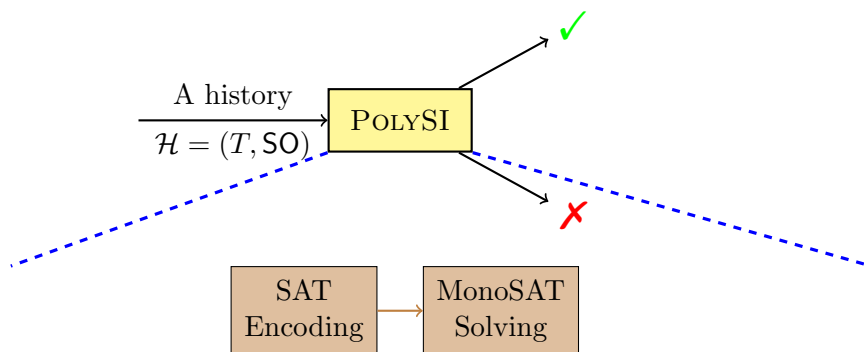


# Contributions: the POLYSI Checker



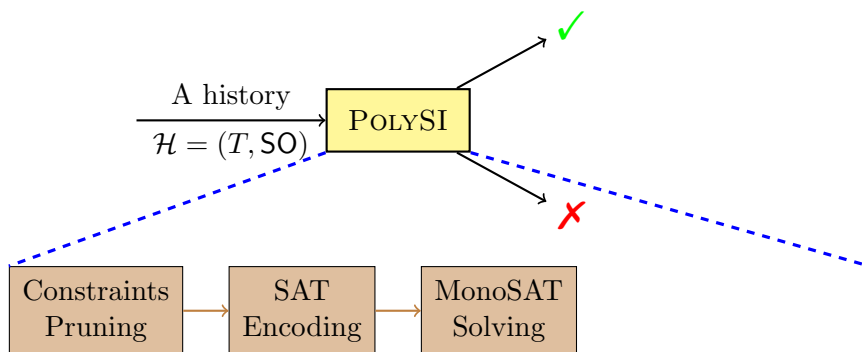
*Sound & Complete:* polygraph-based characterization of SI

# Contributions: the POLYSI Checker



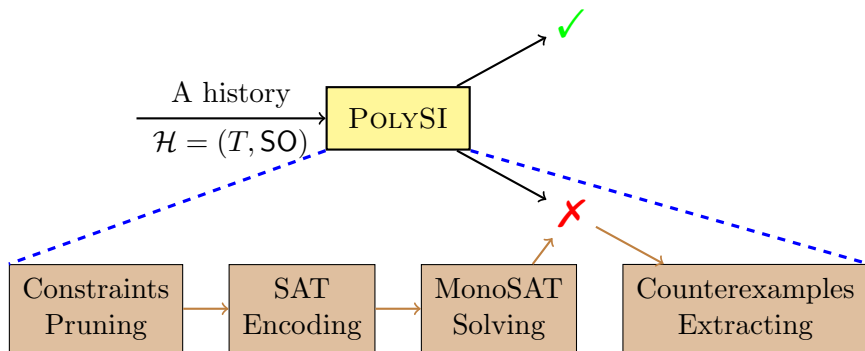
*Efficient:* utilizing MonoSAT solver optimized for graph problems

# Contributions: the POLYSI Checker



*Efficient:* domain-specific pruning before encoding

# Contributions: the POLYSI Checker




*Informative:* extract counterexamples from the unsatisfiable core

# Contributions: PolySI

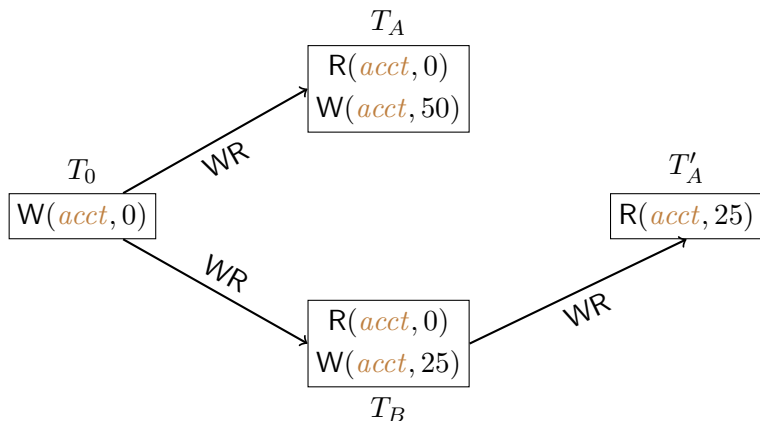
PolySI found SI violations in production database systems.

PolySI outperformed state-of-the-art black-box SI checkers and scales up to large workloads.

# Dependency Graph-based Characterization of SI



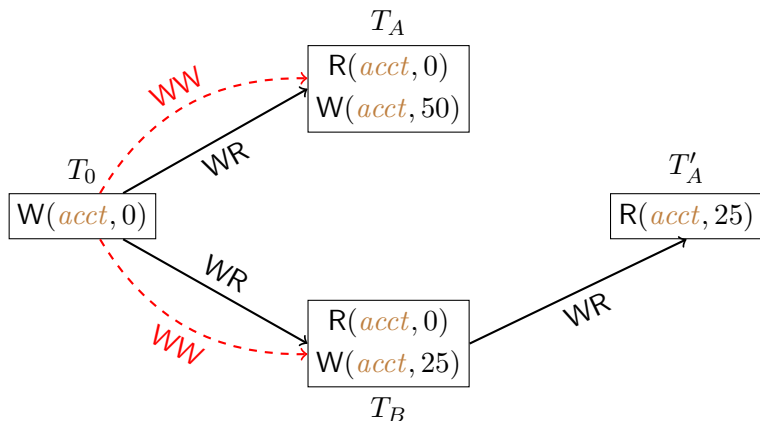
# Dependency Graph-based Characterization of SI



WR: “write-read” dependency capturing the “read-from” relation



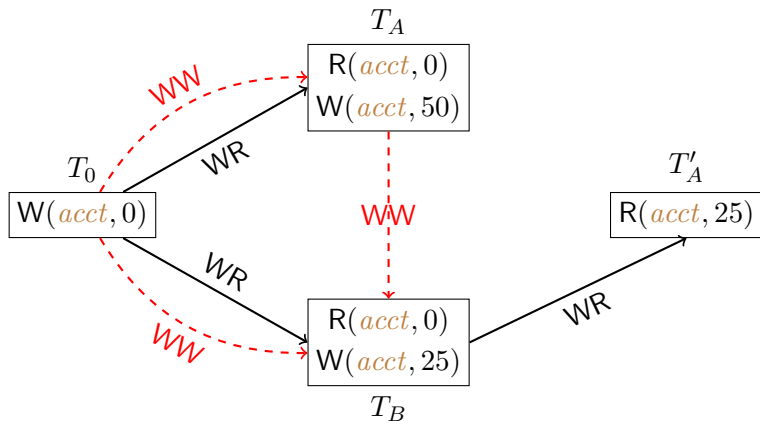
# Dependency Graph-based Characterization of SI



WW: “write-write” dependency capturing the version order

# Dependency Graph-based Characterization of SI

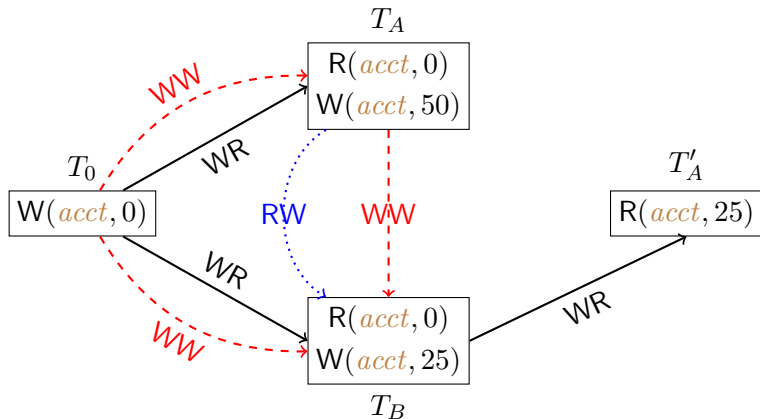
Suppose that  $T_A \xrightarrow{WW} T_B$



WW: “write-write” dependency capturing the version order

# Dependency Graph-based Characterization of SI

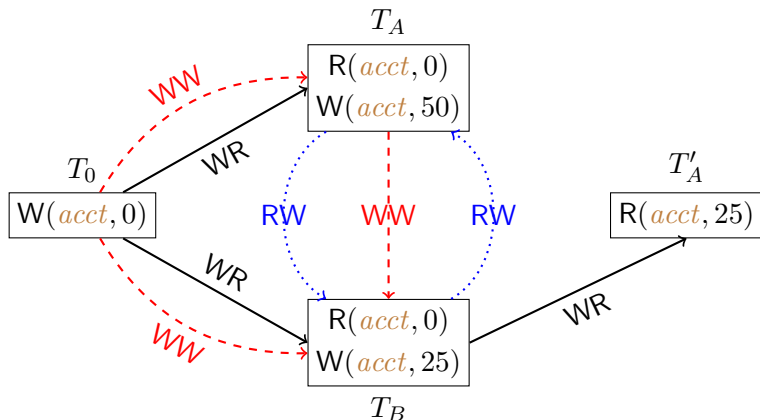
$$T_0 \xrightarrow{WR} T_A \wedge T_0 \xrightarrow{WW} T_B \implies T_A \xrightarrow{RW} T_B$$



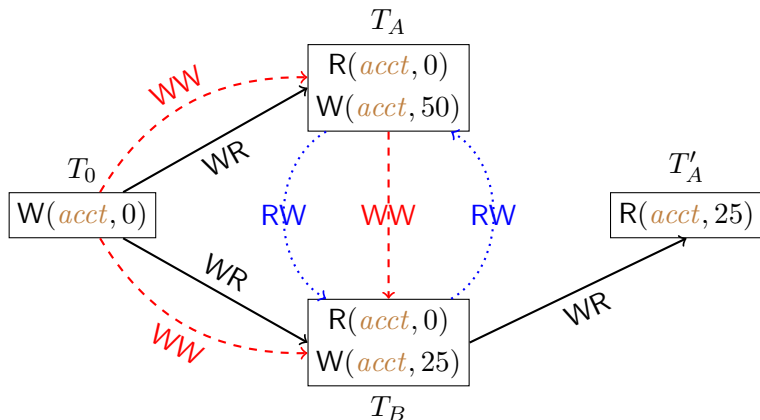
RW: “read-write” dependency capturing the overwritten relation

# Dependency Graph-based Characterization of SI

$$T_0 \xrightarrow{WR} T_B \wedge T_0 \xrightarrow{WW} T_A \implies T_A \xrightarrow{RW} T_A$$



# Dependency Graph-based Characterization of SI



undesiable cycle:  $T_A \xrightarrow{WW} T_B \xrightarrow{RW} T_A$

# Dependency Graph-based Characterization of SI

# Dependency Graph-based Characterization of SI

# Dependency Graph-based Characterization of SI

SI is characterised by dependency graphs that contain only cycles with *at least two adjacent anti-dependency* edges.

Theorem (Theorem 4.1 of [Cerone and Gotsman, 2018])

For a history  $\mathcal{H} = (T, \text{SO})$ ,

$$\begin{aligned} \mathcal{H} \models \text{SI} \iff \mathcal{H} \models \text{INT} \wedge \\ \exists \text{WR, WW, RW. } \mathcal{G} = (\mathcal{H}, \text{WR, WW, RW}) \wedge \\ (((\text{SO}_{\mathcal{G}} \cup \text{WR}_{\mathcal{G}} \cup \text{WW}_{\mathcal{G}}) ; \text{RW}_{\mathcal{G}}?) \text{ is acyclic}). \end{aligned}$$



# Dependency Graph-based Characterization of SI

SI is characterised by dependency graphs that contain only cycles with *at least two adjacent anti-dependency* edges.

Theorem (Theorem 4.1 of [Cerone and Gotsman, 2018])

For a history  $\mathcal{H} = (T, \text{SO})$ ,

$$\begin{aligned} \mathcal{H} \models \text{SI} &\iff \mathcal{H} \models \text{INT} \wedge \\ &\exists \text{WR, WW, RW. } \mathcal{G} = (\mathcal{H}, \text{WR, WW, RW}) \wedge \\ &\quad (((\text{SO}_{\mathcal{G}} \cup \text{WR}_{\mathcal{G}} \cup \text{WW}_{\mathcal{G}}) ; \text{RW}_{\mathcal{G}}?) \text{ is acyclic}). \end{aligned}$$


enumerate all possible **WW** relations

and check whether any of them satisfies the “cycle” condition.



# Dependency Graph-based Characterization of SI

explain ;




# Polygraphs: A Family of Dependency Graphs







# Polygraphs: A Family of Dependency Graphs



# Polygraphs: A Family of Dependency Graphs






# Polygraphs: A Family of Dependency Graphs





polygraph:  $\langle \text{either} \triangleq T_k \xrightarrow{WW} T_i, \text{ or } \triangleq T_j \xrightarrow{RW} T_k \rangle$

# Polygraphs: A Family of Dependency Graphs





# Polygraphs: A Family of Dependency Graphs








# Polygraphs: A Family of Dependency Graphs



# Polygraphs: A Family of Dependency Graphs



generalized polygraph:

$$\langle \text{either} \triangleq \{T \xrightarrow{WW} S, T' \xrightarrow{RW} S\}, \text{ or } \triangleq \{S \xrightarrow{WW} T, S' \xrightarrow{RW} T'\} \rangle$$

$$T_1$$

$$\boxed{W(x, 1)}$$

$$T_3$$

$$\boxed{R(x, 1) R(y, 0)}$$

$$T_0$$

$$\boxed{W(x, 0) W(y, 0)}$$

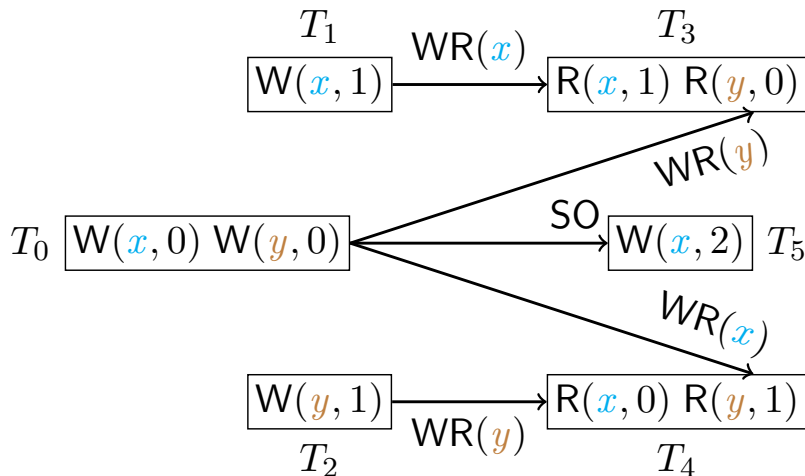
$$\boxed{W(x, 2)} T_5$$

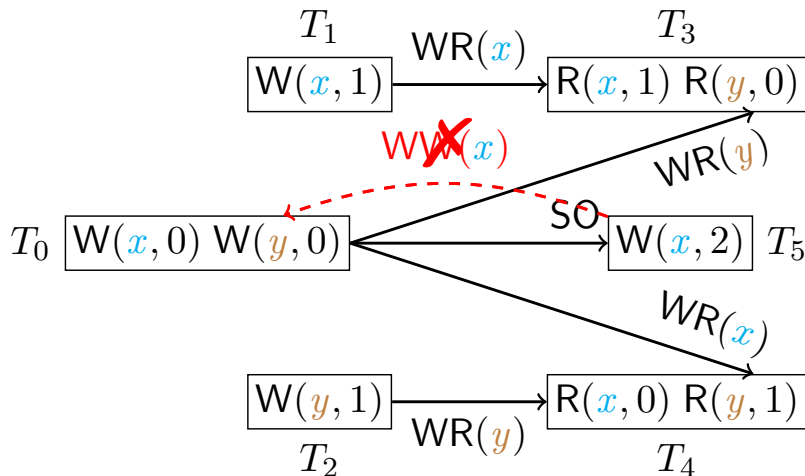
$$\boxed{W(y, 1)}$$

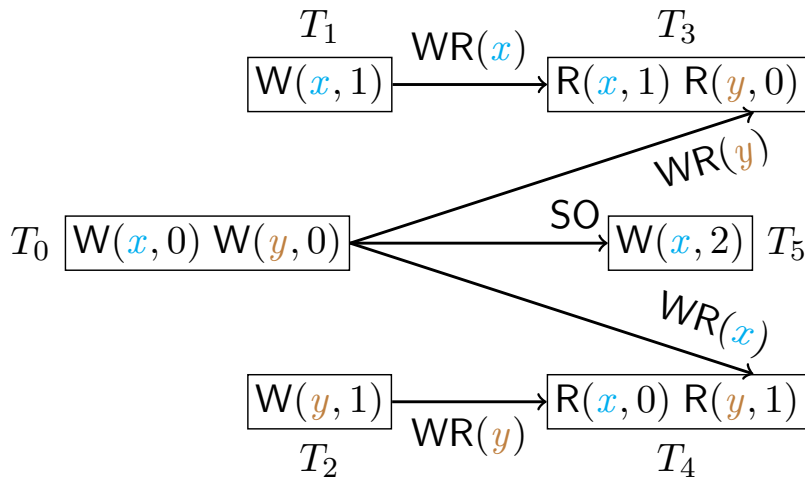
$$T_2$$

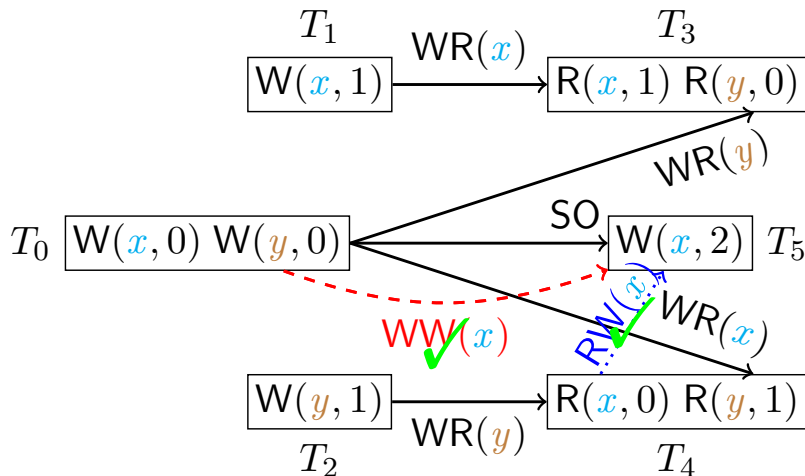
$$\boxed{R(x, 0) R(y, 1)}$$


$$T_4$$



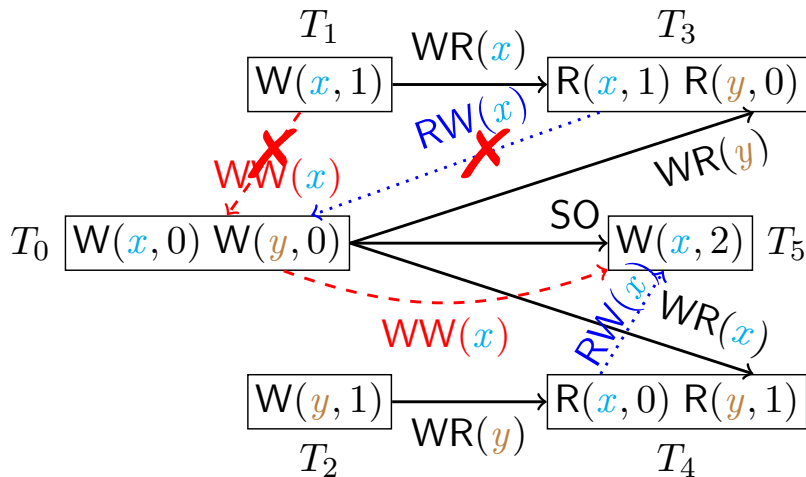


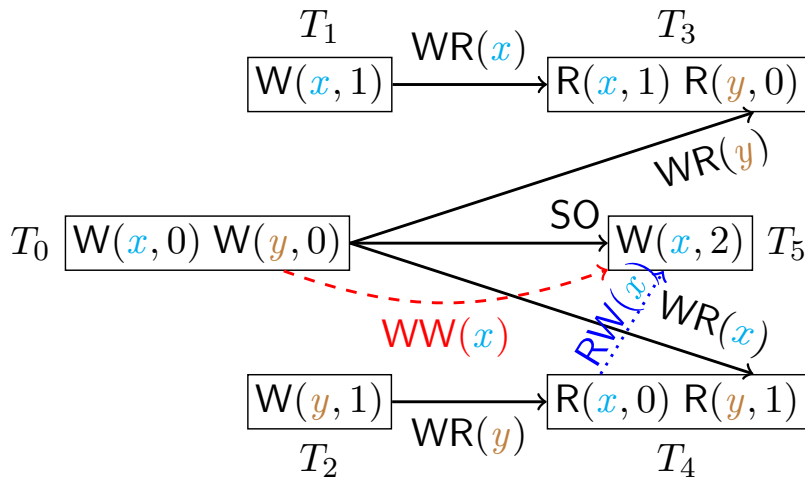





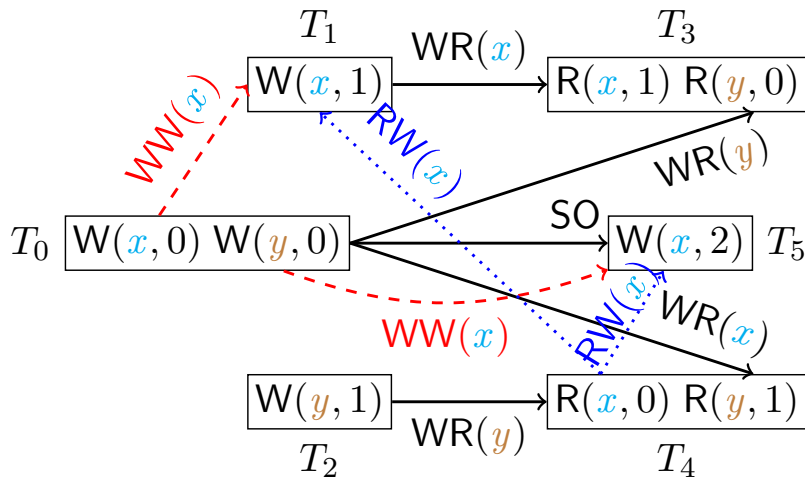


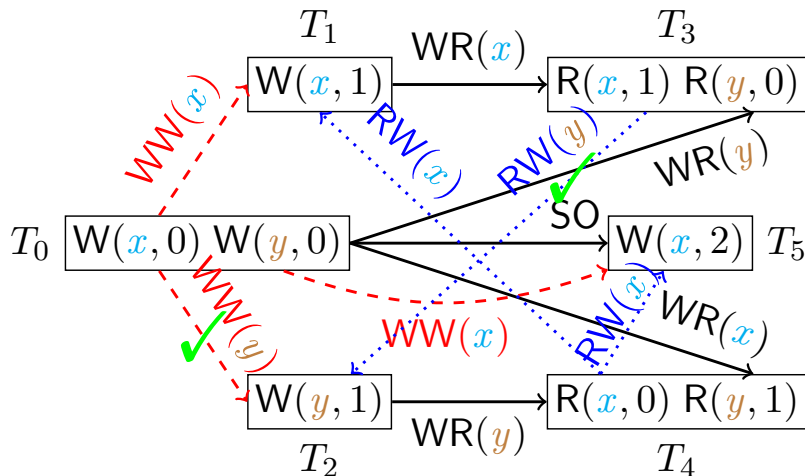



















Hengfeng Wei (hfwei@nju.edu.cn)





Cerone, Andrea and Alexey Gotsman (Jan. 2018). “Analysing Snapshot Isolation”. In: *J. ACM* 65.2. ISSN: 0004-5411. DOI: 10.1145/3152396. URL: <https://doi.org/10.1145/3152396>.