Efficient Black-box Checking of Snapshot Isolation in Databases

(Conference VLDB'2024)

Hengfeng Wei

hfwei@nju.edu.cn

August 15, 2023

Transaction and Isolation Level

A transaction is a *group* of operations that is executed atomically.

$x_1 \leftarrow \mathsf{R}(acct_1)$
$x_2 \leftarrow \mathsf{R}(acct_2)$
if $x_1 + x_2 > 100$
  $x_2 \leftarrow x_2 - 100$
  $\mathsf{W}(acct_2, x_2)$

$acct_1 = acct_2 = 60$

Transaction and Isolation Level

A transaction is a *group* of operations that is executed atomically.

$x_1 \leftarrow \mathsf{R}(acct_1)$
$x_2 \leftarrow \mathsf{R}(acct_2)$
if $x_1 + x_2 > 100$
$\quad x_1 \leftarrow x_1 - 100$
$\quad \mathsf{W}(acct_1, x_1)$

$x_1 \leftarrow \mathsf{R}(acct_1)$
$x_2 \leftarrow \mathsf{R}(acct_2)$
if $x_1 + x_2 > 100$
$\quad x_2 \leftarrow x_2 - 100$
$\quad \mathsf{W}(acct_2, x_2)$

$x_1 \leftarrow \mathsf{R}(acct_1)$
$x_2 \leftarrow \mathsf{R}(acct_2)$

$acct_1 = acct_2 = 60$

Transaction and Isolation Level

A transaction is a *group* of operations that is executed atomically.

$x_1 \leftarrow \mathsf{R}(acct_1)$
$x_2 \leftarrow \mathsf{R}(acct_2)$
if $x_1 + x_2 > 100$
  $x_1 \leftarrow x_1 - 100$
  $\mathsf{W}(acct_1, x_1)$

$x_1 \leftarrow \mathsf{R}(acct_1)$
$x_2 \leftarrow \mathsf{R}(acct_2)$
if $x_1 + x_2 > 100$
  $x_2 \leftarrow x_2 - 100$
  $\mathsf{W}(acct_2, x_2)$

$x_1 \leftarrow \mathsf{R}(acct_1)$
$x_2 \leftarrow \mathsf{R}(acct_2)$

Isolation Level (aka Transactional Consistency Model)

$acct_1 = acct_2 = 60$

The isolation levels specify how they are isolated from each other.

Serializability (SER)

All transactions appear to execute in some total order.

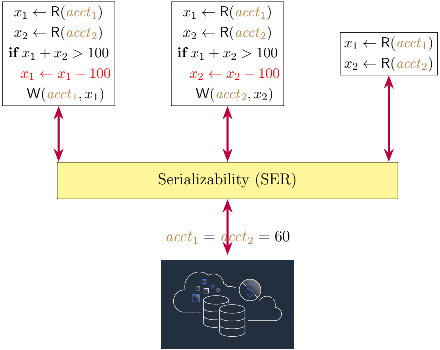

Serializability (SER)

All transactions appear to execute in some total order.

$x_1 \leftarrow \mathsf{R}(acct_1)$
$x_2 \leftarrow \mathsf{R}(acct_2)$
if $x_1 + x_2 > 100$
$x_1 \leftarrow x_1 - 100$
$\mathsf{W}(acct_1, x_1)$

$x_1 \leftarrow \mathsf{R}(acct_1)$
$x_2 \leftarrow \mathsf{R}(acct_2)$
if $x_1 + x_2 > 100$
$x_2 \leftarrow x_2 - 100$
$\mathsf{W}(acct_2, x_2)$

$x_1 \leftarrow \mathsf{R}(acct_1)$
$x_2 \leftarrow \mathsf{R}(acct_2)$

Serializability (SER)

$acct_1 = acct_2 = 60$

$\mathsf{R}(acct_1, 60)$
$\mathsf{R}(acct_2, 60)$
$\mathsf{W}(acct_1, -40)$

Serializability (SER)

All transactions appear to execute in some total order.

$x_1 \leftarrow \mathsf{R}(acct_1)$
$x_2 \leftarrow \mathsf{R}(acct_2)$
if $x_1 + x_2 > 100$
$x_1 \leftarrow x_1 - 100$
$\mathsf{W}(acct_1, x_1)$

$x_1 \leftarrow \mathsf{R}(acct_1)$
$x_2 \leftarrow \mathsf{R}(acct_2)$
if $x_1 + x_2 > 100$
$x_2 \leftarrow x_2 - 100$
$\mathsf{W}(acct_2, x_2)$

$x_1 \leftarrow \mathsf{R}(acct_1)$
$x_2 \leftarrow \mathsf{R}(acct_2)$

Serializability (SER)

$acct_1 = acct_2 = 60$

$\mathsf{R}(acct_1, 60)$
$\mathsf{R}(acct_2, 60)$
$\mathsf{W}(acct_1, -40)$

$\mathsf{R}(acct_1, 60)$
$\mathsf{R}(acct_2, 60)$

Serializability (SER)

All transactions appear to execute in some total order.

$x_1 \leftarrow \mathsf{R}(acct_1)$
$x_2 \leftarrow \mathsf{R}(acct_2)$
if $x_1 + x_2 > 100$
$x_1 \leftarrow x_1 - 100$
$\mathsf{W}(acct_1, x_1)$

$x_1 \leftarrow \mathsf{R}(acct_1)$
$x_2 \leftarrow \mathsf{R}(acct_2)$
if $x_1 + x_2 > 100$
$x_2 \leftarrow x_2 - 100$
$\mathsf{W}(acct_2, x_2)$

$x_1 \leftarrow \mathsf{R}(acct_1)$
$x_2 \leftarrow \mathsf{R}(acct_2)$

Serializability (SER)

$acct_1 = acct_2 = 60$

$\mathsf{R}(acct_1, 60)$
$\mathsf{R}(acct_2, 60)$
$\mathsf{W}(acct_1, -40)$

$\mathsf{R}(acct_1, 60)$
$\mathsf{R}(acct_2, 60)$

$\mathsf{R}(acct_1, 60)$
$\mathsf{R}(acct_2, 60)$

Serializability (SER)

All transactions appear to execute in some total order.

$x_1 \leftarrow \mathsf{R}(acct_1)$
$x_2 \leftarrow \mathsf{R}(acct_2)$
if $x_1 + x_2 > 100$
$x_1 \leftarrow x_1 - 100$
$\mathsf{W}(acct_1, x_1)$

$x_1 \leftarrow \mathsf{R}(acct_1)$
$x_2 \leftarrow \mathsf{R}(acct_2)$
if $x_1 + x_2 > 100$
$x_2 \leftarrow x_2 - 100$
$\mathsf{W}(acct_2, x_2)$

$x_1 \leftarrow \mathsf{R}(acct_1)$
$x_2 \leftarrow \mathsf{R}(acct_2)$

Serializability (SER)

$acct_1 = acct_2 = 60$

$\mathsf{R}(acct_1, 60)$
$\mathsf{R}(acct_2, 60)$
$\mathsf{W}(acct_1, -40)$

$\mathsf{R}(acct_1, 60)$
$\mathsf{R}(acct_2, 60)$

$\mathsf{R}(acct_1, 60)$
$\mathsf{R}(acct_2, 60)$

too expensive, especially for distributed transactions

Snapshot Isolation (SI)

$x_1 \leftarrow \mathsf{R}(acct_1)$
$x_2 \leftarrow \mathsf{R}(acct_2)$
if $x_1 + x_2 > 100$
$x_1 \leftarrow x_1 - 100$
$\mathsf{W}(acct_1, x_1)$

$x_1 \leftarrow \mathsf{R}(acct_1)$
$x_2 \leftarrow \mathsf{R}(acct_2)$
if $x_1 + x_2 > 100$
$x_2 \leftarrow x_2 - 100$
$\mathsf{W}(acct_2, x_2)$

$x_1 \leftarrow \mathsf{R}(acct_1)$
$x_2 \leftarrow \mathsf{R}(acct_2)$

Snapshot Isolation (SI)

$acct_1 = acct_2 = 60$

Snapshot Isolation (SI)

$x_1 \leftarrow \mathsf{R}(acct_1)$
$x_2 \leftarrow \mathsf{R}(acct_2)$
if $x_1 + x_2 > 100$
$x_1 \leftarrow x_1 - 100$
$\mathsf{W}(acct_1, x_1)$

$x_1 \leftarrow \mathsf{R}(acct_1)$
$x_2 \leftarrow \mathsf{R}(acct_2)$
if $x_1 + x_2 > 100$
$x_2 \leftarrow x_2 - 100$
$\mathsf{W}(acct_2, x_2)$

$x_1 \leftarrow \mathsf{R}(acct_1)$
$x_2 \leftarrow \mathsf{R}(acct_2)$

Snapshot Isolation (SI)

$acct_1 = acct_2 = 60$

$\mathsf{R}(acct_1, 60)$
$\mathsf{R}(acct_2, 60)$
$\mathsf{W}(acct_1, -40)$

Snapshot Isolation (SI)

$x_1 \leftarrow \mathsf{R}(acct_1)$
$x_2 \leftarrow \mathsf{R}(acct_2)$
if $x_1 + x_2 > 100$
$x_1 \leftarrow x_1 - 100$
$\mathsf{W}(acct_1, x_1)$

$x_1 \leftarrow \mathsf{R}(acct_1)$
$x_2 \leftarrow \mathsf{R}(acct_2)$
if $x_1 + x_2 > 100$
$x_2 \leftarrow x_2 - 100$
$\mathsf{W}(acct_2, x_2)$

$x_1 \leftarrow \mathsf{R}(acct_1)$
$x_2 \leftarrow \mathsf{R}(acct_2)$

Snapshot Isolation (SI)

$acct_1 = acct_2 = 60$

$\mathsf{R}(acct_1, 60)$
$\mathsf{R}(acct_2, 60)$
$\mathsf{W}(acct_1, -40)$

$\mathsf{R}(acct_1, 60)$
$\mathsf{R}(acct_2, 60)$
$\mathsf{W}(acct_2, -40)$

Snapshot Isolation (SI)

$x_1 \leftarrow \mathsf{R}(acct_1)$
$x_2 \leftarrow \mathsf{R}(acct_2)$
if $x_1 + x_2 > 100$
$x_1 \leftarrow x_1 - 100$
$\mathsf{W}(acct_1, x_1)$

$x_1 \leftarrow \mathsf{R}(acct_1)$
$x_2 \leftarrow \mathsf{R}(acct_2)$
if $x_1 + x_2 > 100$
$x_2 \leftarrow x_2 - 100$
$\mathsf{W}(acct_2, x_2)$

$x_1 \leftarrow \mathsf{R}(acct_1)$
$x_2 \leftarrow \mathsf{R}(acct_2)$

Snapshot Isolation (SI)

$acct_1 = acct_2 = 60$

$\mathsf{R}(acct_1, 60)$
$\mathsf{R}(acct_2, 60)$
$\mathsf{W}(acct_1, -40)$

$\mathsf{R}(acct_1, 60)$
$\mathsf{R}(acct_2, 60)$
$\mathsf{W}(acct_2, -40)$

$\mathsf{R}(acct_1, -40)$
$\mathsf{R}(acct_2, -40)$

Snapshot Isolation (SI)

$x_1 \leftarrow \mathsf{R}(acct_1)$
$x_2 \leftarrow \mathsf{R}(acct_2)$
if $x_1 + x_2 > 100$
$x_1 \leftarrow x_1 - 100$
$\mathsf{W}(acct_1, x_1)$

$x_1 \leftarrow \mathsf{R}(acct_1)$
$x_2 \leftarrow \mathsf{R}(acct_2)$
if $x_1 + x_2 > 100$
$x_2 \leftarrow x_2 - 100$
$\mathsf{W}(acct_2, x_2)$

$x_1 \leftarrow \mathsf{R}(acct_1)$
$x_2 \leftarrow \mathsf{R}(acct_2)$

Snapshot Isolation (SI)

$acct_1 = acct_2 = 60$

Write Skew

$\mathsf{R}(acct_1, 60)$
$\mathsf{R}(acct_2, 60)$
$\mathsf{W}(acct_1, -40)$

$\mathsf{R}(acct_1, 60)$
$\mathsf{R}(acct_2, 60)$
$\mathsf{W}(acct_2, -40)$

$\mathsf{R}(acct_1, -40)$
$\mathsf{R}(acct_2, -40)$

Snapshot Isolation (SI)

$x_1 \leftarrow \mathsf{R}(acct_1)$
$x_2 \leftarrow \mathsf{R}(acct_2)$
if $x_1 + x_2 > 100$
$x_1 \leftarrow x_1 - 100$
$\mathsf{W}(acct_1, x_1)$

$x_1 \leftarrow \mathsf{R}(acct_1)$
$x_2 \leftarrow \mathsf{R}(acct_2)$
if $x_1 + x_2 > 100$
$x_2 \leftarrow x_2 - 100$
$\mathsf{W}(acct_2, x_2)$

$x_1 \leftarrow \mathsf{R}(acct_1)$
$x_2 \leftarrow \mathsf{R}(acct_2)$

Snapshot Isolation (SI)

$acct_1 = acct_2 = 60$

Write Skew ✓

$\mathsf{R}(acct_1, 60)$
$\mathsf{R}(acct_2, 60)$
$\mathsf{W}(acct_1, -40)$

$\mathsf{R}(acct_1, 60)$
$\mathsf{R}(acct_2, 60)$
$\mathsf{W}(acct_2, -40)$

$\mathsf{R}(acct_1, -40)$
$\mathsf{R}(acct_2, -40)$

Snapshot Isolation (SI)

$x_1 \leftarrow \mathsf{R}(acct_1)$
$x_2 \leftarrow \mathsf{R}(acct_2)$
if $x_1 + x_2 > 100$
$x_1 \leftarrow x_1 - 100$
$\mathsf{W}(acct_1, x_1)$

$x_1 \leftarrow \mathsf{R}(acct_1)$
$x_2 \leftarrow \mathsf{R}(acct_2)$
if $x_1 + x_2 > 100$
$x_2 \leftarrow x_2 - 100$
$\mathsf{W}(acct_2, x_2)$

$x_1 \leftarrow \mathsf{R}(acct_1)$
$x_2 \leftarrow \mathsf{R}(acct_2)$

Snapshot Isolation (SI)

$acct_1 = acct_2 = 60$

WRITE SKEW ✓

$\mathsf{R}(acct_1, 60)$
$\mathsf{R}(acct_2, 60)$
$\mathsf{W}(acct_1, -40)$

$\mathsf{R}(acct_1, 60)$
$\mathsf{R}(acct_2, 60)$
$\mathsf{W}(acct_2, -40)$

$\mathsf{R}(acct_1, -40)$
$\mathsf{R}(acct_2, -40)$

Snapshot Read: Each transaction reads data from a *snapshot* of committed data valid as of the (logical) time the transaction started.

Snapshot Isolation (SI)

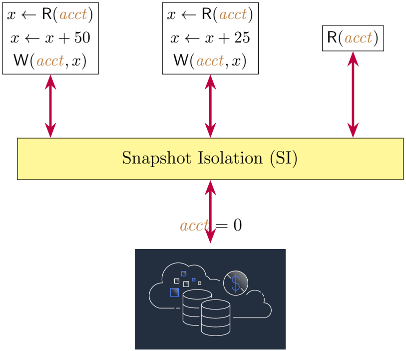

Snapshot Isolation (SI)

$x \leftarrow \mathsf{R}(acct)$
$x \leftarrow x + 50$
$\mathsf{W}(acct, x)$

$x \leftarrow \mathsf{R}(acct)$
$x \leftarrow x + 25$
$\mathsf{W}(acct, x)$

$\mathsf{R}(acct)$

Snapshot Isolation (SI)

$acct = 0$

$\mathsf{R}(acct, 0)$
$\mathsf{W}(acct, 50)$

Snapshot Isolation (SI)

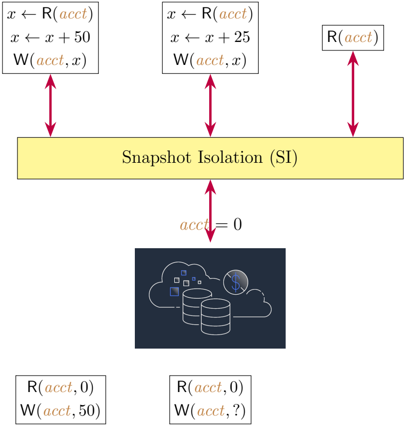

Snapshot Isolation (SI)

$x \leftarrow \mathsf{R}(acct)$
$x \leftarrow x + 50$
$\mathsf{W}(acct, x)$

$x \leftarrow \mathsf{R}(acct)$
$x \leftarrow x + 25$
$\mathsf{W}(acct, x)$

$\mathsf{R}(acct)$

Snapshot Isolation (SI)

$acct = 0$

$\mathsf{R}(acct, 0)$
$\mathsf{W}(acct, 50)$

$\mathsf{R}(acct, 0)$
$\mathsf{W}(acct, ?)$

Snapshot Write: Concurrent transactions cannot write to the same key. One of them must be aborted.

Snapshot Isolation (SI)

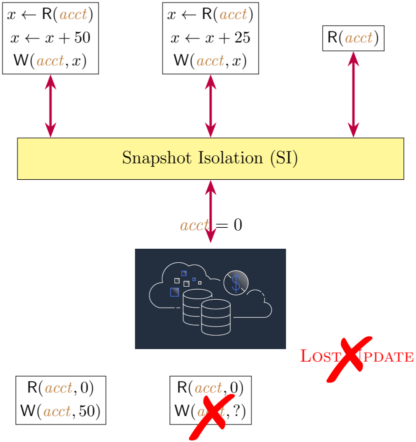

Snapshot Write: Concurrent transactions cannot write to the same key. One of them must be aborted.

Snapshot Isolation (SI)

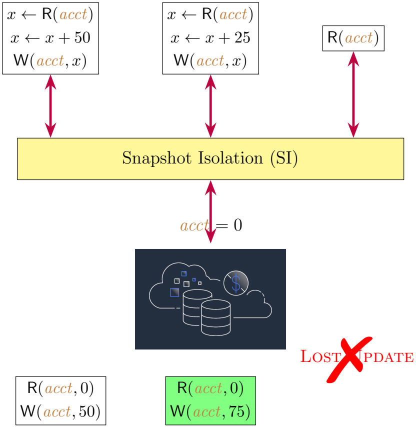

Snapshot Write: Concurrent transactions cannot write to the same key. One of them must be aborted.

Snapshot Isolation (SI)

Snapshot Isolation (SI)

Snapshot Isolation (SI)

Snapshot Isolation (SI)

$\boxed{\mathsf{W}(x, post)}$

Snapshot Isolation (SI)

Snapshot Isolation (SI)

$\mathsf{W}(x, post)$

$\mathsf{R}(x, post)$
$\mathsf{W}(y, comment)$

Snapshot Isolation (SI)

Snapshot Isolation (SI)

W(x, *post*)

R(x, *post*)
W(y, *comment*)

R(x, *empty*)
R(y, *comment*)

Snapshot Isolation (SI)

Snapshot Isolation (SI)

CAUSALITY VIOLATION

W(x, *post*)	R(x, *post*) W(y, *comment*)	R(x, *empty*) R(y, *comment*)

Snapshot Isolation (SI)

Snapshot Isolation (SI)

CAUSALITY VIOLATION ✗

$\mathsf{W}(x, post)$	$\mathsf{R}(x, post)$ $\mathsf{W}(y, comment)$	$\mathsf{R}(x, empty)$ $\mathsf{R}(y, comment)$ ✗

Databases and Snapshot Isolation

database logos

Many databases claim to support SI.

Databases and Snapshot Isolation

+papers

Databases may fail to provide SI as they claim.

The SI Checking Problem

Definition (The SI Checking Problem)

The SI checking problem is the decision problem of determing whether a given history $\mathcal{H} = (T, \mathsf{SO})$ satisfies SI?

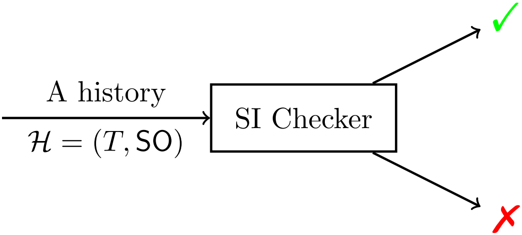

The SI Checking Problem

Definition (The SI Checking Problem)

The SI checking problem is the decision problem of determing whether a given history $\mathcal{H} = (T, \mathsf{SO})$ satisfies SI?

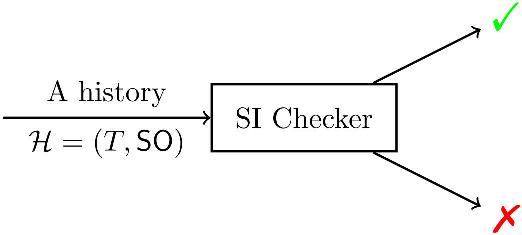

SO : *session order* among the set T of transactions

The SI Checking Problem

Black-box checking: do not rely on database internals

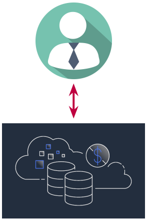

The histories are collected from database logs.

The SI Checking Problem

Black-box checking: do not rely on database internals

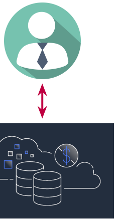

$\mathsf{W}(x, \textit{post})$

$\mathsf{R}(x, \textit{post})$
$\mathsf{W}(y, \textit{comment})$

$\mathsf{R}(x, \textit{empty})$
$\mathsf{R}(y, \textit{comment})$

Lost Update

The histories are collected from database logs.

The SI Checking Problem

Black-box checking: do not rely on database internals

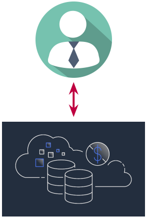

W(x, *post*)	R(x, *post*) W(y, *comment*)	R(x, *empty*) R(y, *comment*)	Lost Update ✗

R($acct_1$, 60) R($acct_2$, 60) W($acct_1$, −40)	R($acct_1$, 60) R($acct_2$, 60) W($acct_2$, −40)	R($acct_1$, −40) R($acct_2$, −40)	Write Skew ✓

The histories are collected from database logs.

The SI Checking Problem

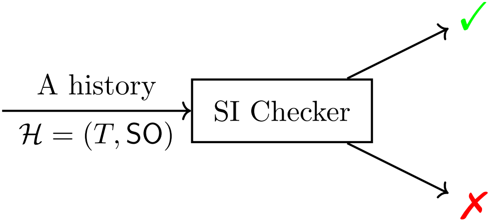

The SI Checking Problem

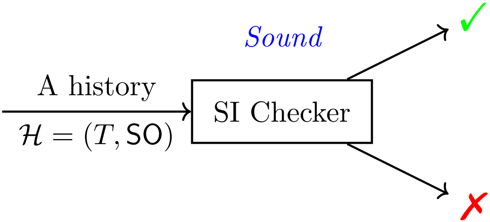

Sound: If the checker says ✗, then the history does *not* satisfy SI.

The SI Checking Problem

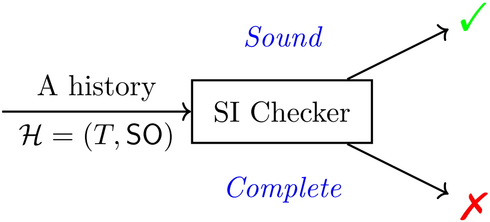

Complete: If the checker says ✓, then the history *satisfies* SI.

The SI Checking Problem

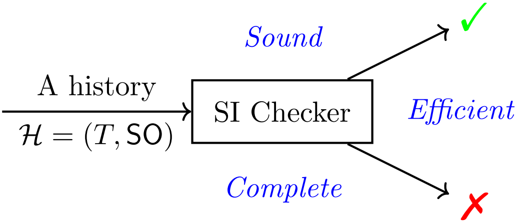

Efficient: The checker should *scale* up to large workloads.

The SI Checking Problem

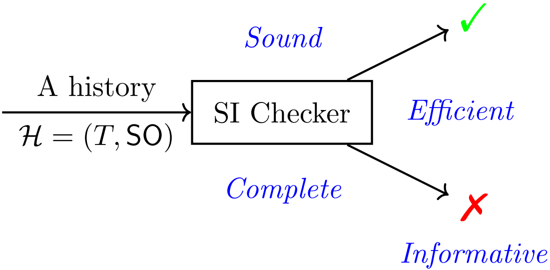

Informative: The checker should provide understandable *counterexamples* if it says ✗.

The SI Checking Problem

related-work

Contribution: the POLYSI Checker

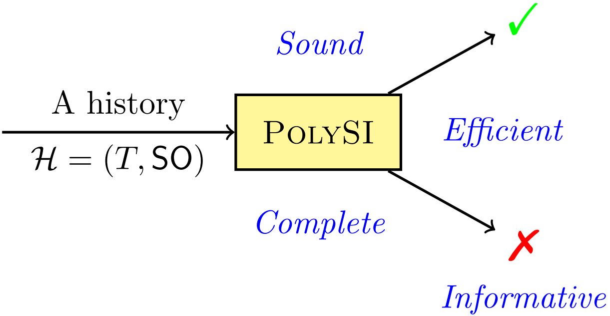

Contribution: the PolySI Checker

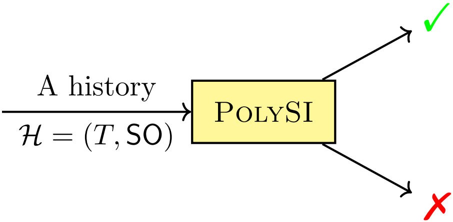

Contribution: the PolySI Checker

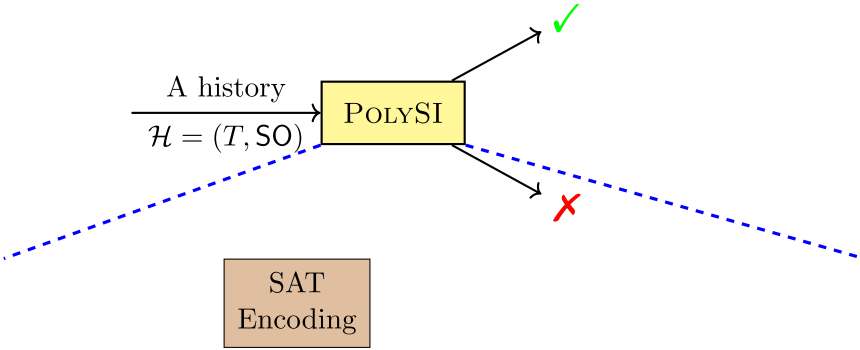

Sound & Complete: polygraph-based characterization of SI

Contribution: the PolySI Checker

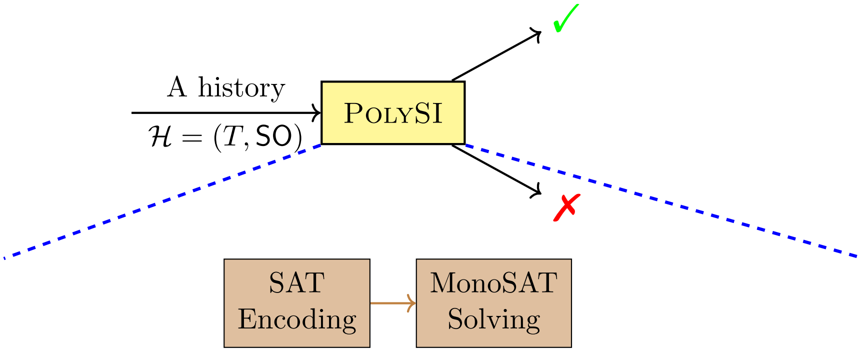

Efficient: utilizing MonoSAT solver optimized for graph problems

Contribution: the PolySI Checker

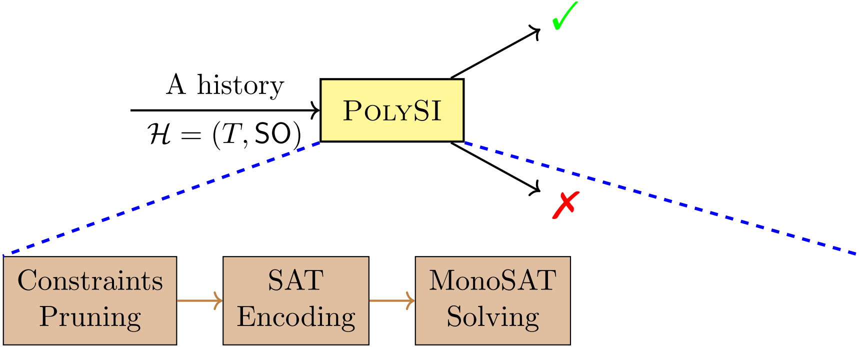

Efficient: domain-specific pruning before encoding

Contribution: the PolySI Checker

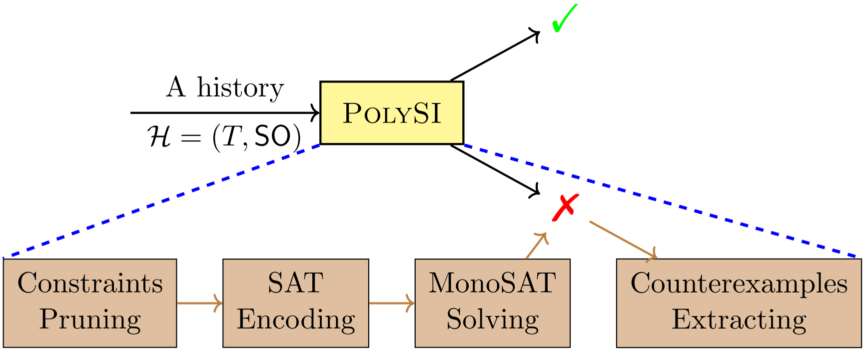

Informative: extract counterexamples from the unsatisifiable core

Dependency Graph-based Characterization of SI

T_A

$\mathsf{R}(acct, 0)$
$\mathsf{W}(acct, 50)$

T_0

$\mathsf{W}(acct, 0)$

T'_A

$\mathsf{R}(acct, 25)$

$\mathsf{R}(acct, 0)$
$\mathsf{W}(acct, 25)$

T_B

Dependency Graph-based Characterization of SI

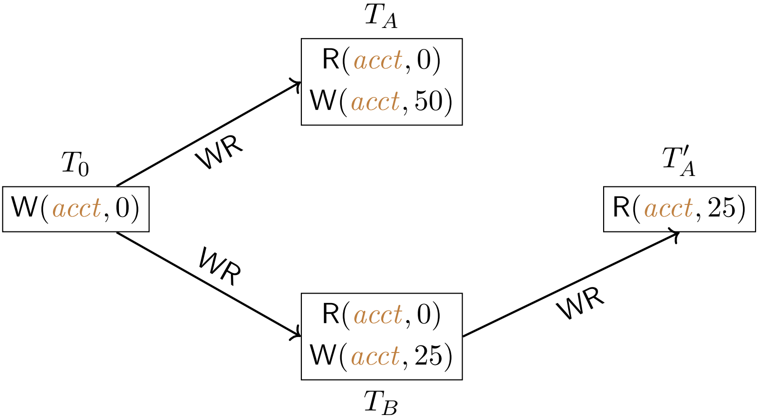

WR: "write-read" dependency capturing the "read-from" relation

Dependency Graph-based Characterization of SI

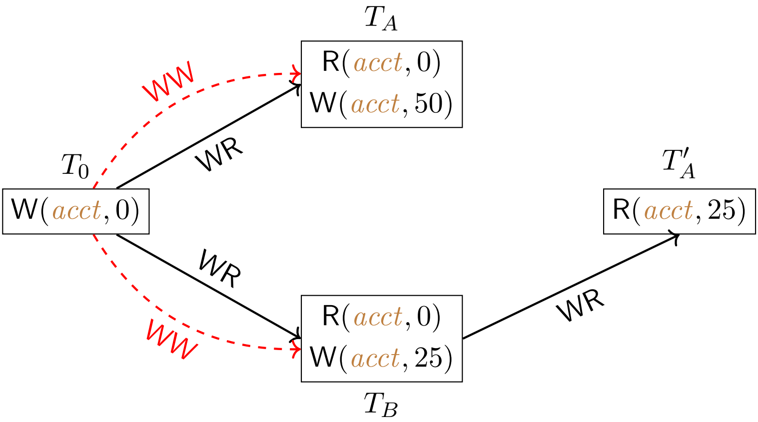

WW: "write-write" dependency capturing the version order

Dependency Graph-based Characterization of SI

Suppose that $T_A \xrightarrow{\text{WW}} T_B$

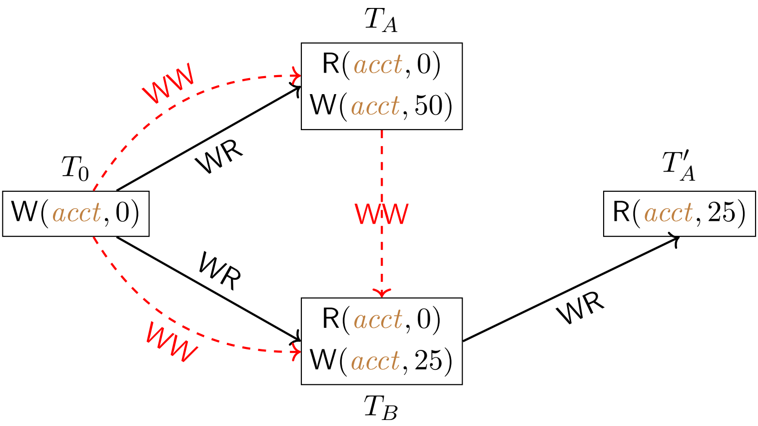

WW: "write-write" dependency capturing the version order

Dependency Graph-based Characterization of SI

$$T_0 \xrightarrow{\mathsf{WR}} T_A \wedge T_0 \xrightarrow{\mathsf{WW}} T_B \implies T_A \xrightarrow{\mathsf{RW}} T_B$$

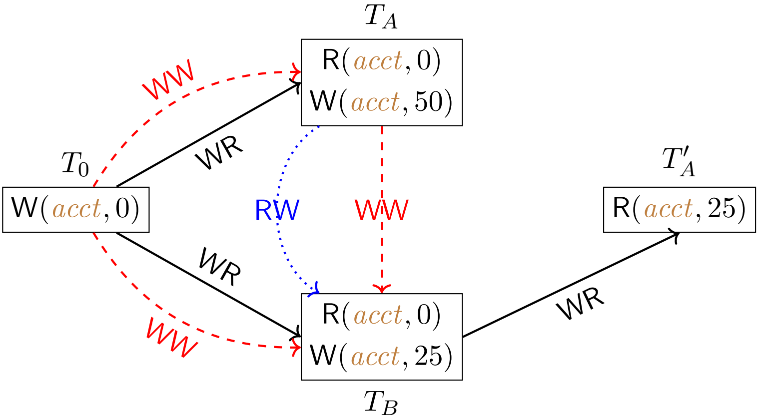

RW: "read-write" dependency capturing the overwritten relation

Dependency Graph-based Characterization of SI

$$T_0 \xrightarrow{\mathsf{WR}} T_B \wedge T_0 \xrightarrow{\mathsf{WW}} T_A \implies T_A \xrightarrow{\mathsf{RW}} T_A$$

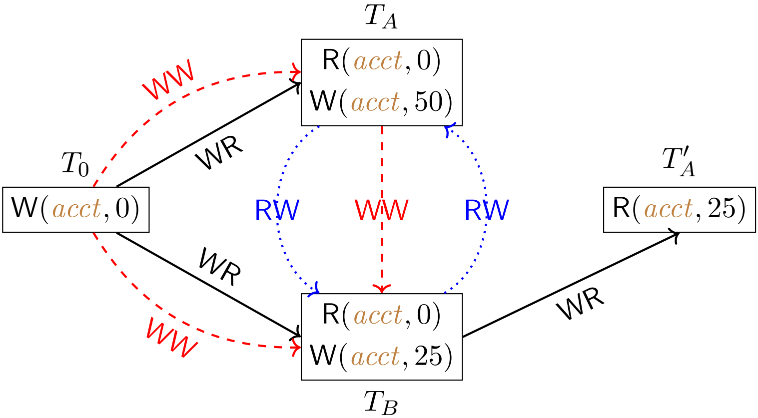

RW: "read-write" dependency capturing the overwritten relation

Dependency Graph-based Characterization of SI

$\boxed{\textit{Suppose that } T_A \xrightarrow{\mathsf{WW}} T_B}$

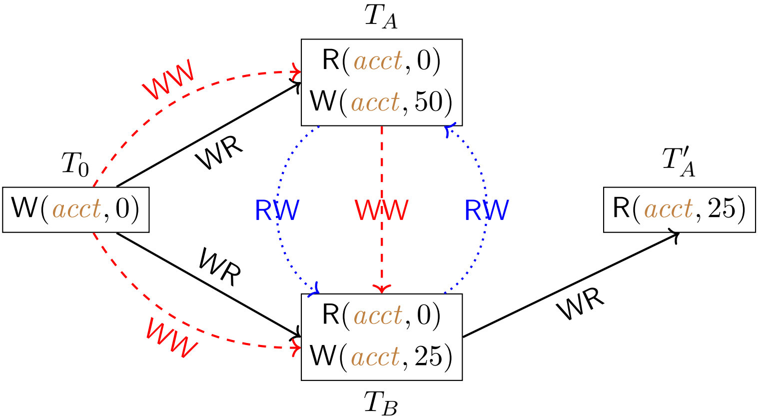

undesired cycle: $T_A \xrightarrow{\mathsf{WW}} T_B \xrightarrow{\mathsf{RW}} T_A$

Dependency Graph-based Characterization of SI

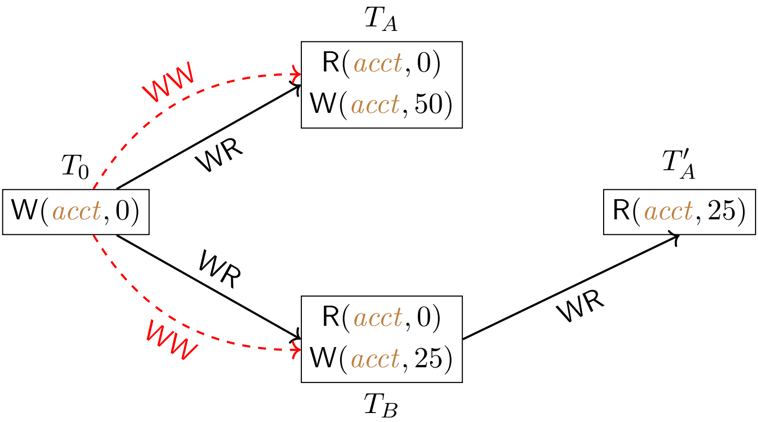

Dependency Graph-based Characterization of SI

Suppose that $T_B \xrightarrow{\text{WW}} T_A$

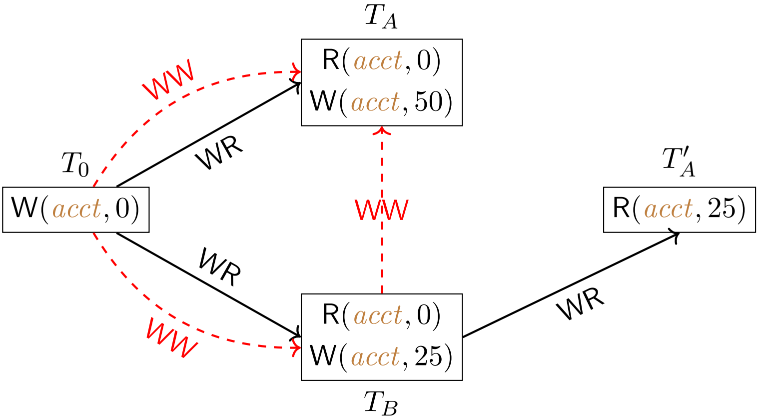

Dependency Graph-based Characterization of SI

Suppose that $T_B \xrightarrow{\textsf{WW}} T_A$

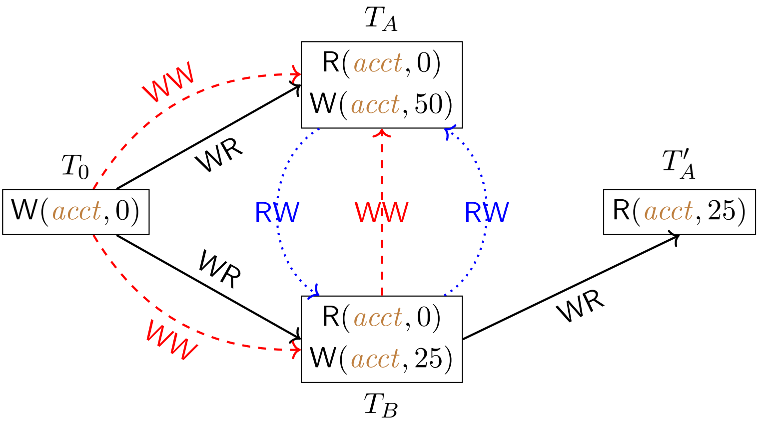

Dependency Graph-based Characterization of SI

Suppose that $T_B \xrightarrow{\mathsf{WW}} T_A$

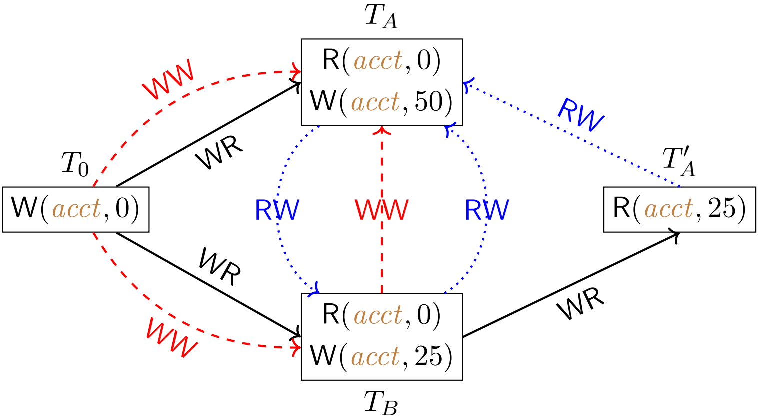

Dependency Graph-based Characterization of SI

Suppose that $T_B \xrightarrow{\text{WW}} T_A$

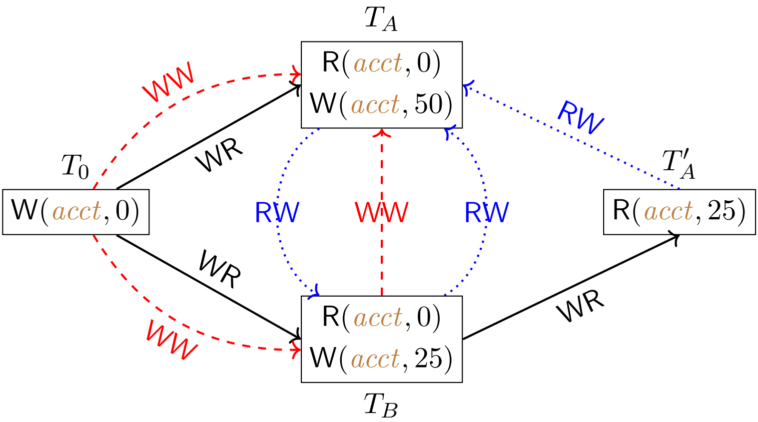

undesired cycle: $T_B \xrightarrow{\text{WW}} T_A \xrightarrow{\text{RW}} T_B$

Dependency Graph-based Characterization of SI

We have considered both bases $T_A \xrightarrow{\mathsf{WW}} T_B$ and $T_B \xrightarrow{\mathsf{WW}} T_A$.

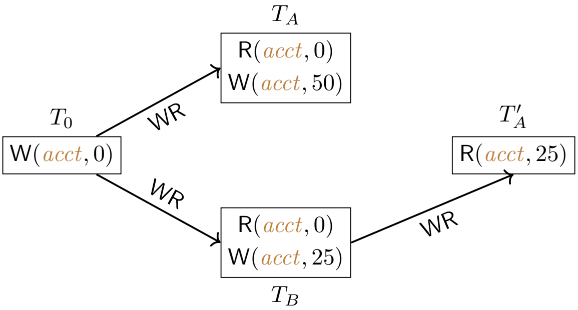

Either case leads to an undesired cycle.

Therefore, it does not satisfy SI.

Dependency Graph-based Characterization of SI

Theorem (Theorem 4.1 of [Cerone and Gotsman, 2018])

Informally, a history satisfies SI if only if there exists a dependency graph for it that contains only cycles (if any) with at least two adjacent RW *edges.*

Dependency Graph-based Characterization of SI

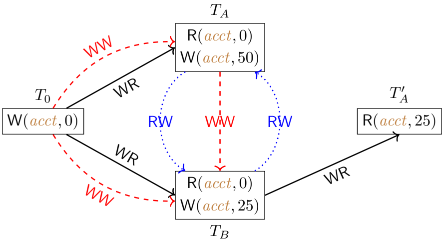

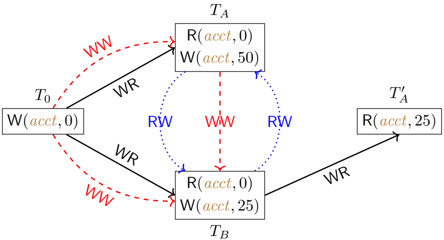

Every possible dependency graph contains an undesired 

cycle.

Dependency Graph-based Characterization of SI

Theorem (Theorem 4.1 of [Cerone and Gotsman, 2018])

For a history $\mathcal{H} = (T, \mathsf{SO})$,

$$\mathcal{H} \models \mathrm{SI} \iff \mathcal{H} \models \textsc{Int} \land \exists\, \mathsf{WR}, \mathsf{WW}, \mathsf{RW}.\ \mathcal{G} = (\mathcal{H}, \mathsf{WR}, \mathsf{WW}, \mathsf{RW}) \land (((\mathsf{SO}_{\mathcal{G}} \cup \mathsf{WR}_{\mathcal{G}} \cup \mathsf{WW}_{\mathcal{G}}) \,;\, \mathsf{RW}_{\mathcal{G}}?) \textit{ is acyclic}).$$

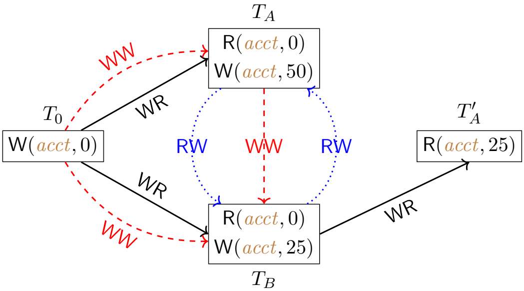

Dependency Graph-based Characterization of SI

Theorem (Theorem 4.1 of [Cerone and Gotsman, 2018])

For a history $\mathcal{H} = (T, \mathsf{SO})$,

$$\mathcal{H} \models \text{SI} \iff \mathcal{H} \models \textsc{Int} \land$$
$$\exists\, \mathsf{WR}, \mathsf{WW}, \mathsf{RW}.\ \mathcal{G} = (\mathcal{H}, \mathsf{WR}, \mathsf{WW}, \mathsf{RW}) \land$$
$$(((\mathsf{SO}_{\mathcal{G}} \cup \mathsf{WR}_{\mathcal{G}} \cup \mathsf{WW}_{\mathcal{G}}) \mathbin{;} \mathsf{RW}_{\mathcal{G}}?) \textit{ is acyclic}).$$

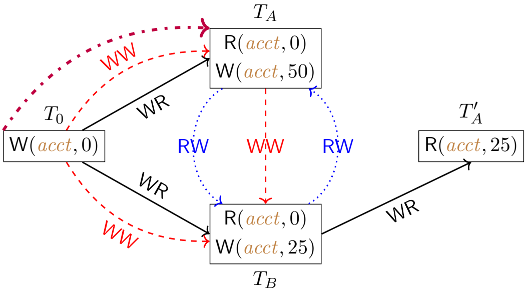

Dependency Graph-based Characterization of SI

Theorem (Theorem 4.1 of [Cerone and Gotsman, 2018])

For a history $\mathcal{H} = (T, \mathsf{SO})$,

$$\mathcal{H} \models \mathrm{SI} \iff \mathcal{H} \models \textsc{Int} \wedge$$
$$\exists\, \mathsf{WR}, \mathsf{WW}, \mathsf{RW}.\ \mathcal{G} = (\mathcal{H}, \mathsf{WR}, \mathsf{WW}, \mathsf{RW}) \wedge$$
$$(((\mathsf{SO}_{\mathcal{G}} \cup \mathsf{WR}_{\mathcal{G}} \cup \mathsf{WW}_{\mathcal{G}}) \ ; \ \mathsf{RW}_{\mathcal{G}}?) \ \textit{ is acyclic}).$$

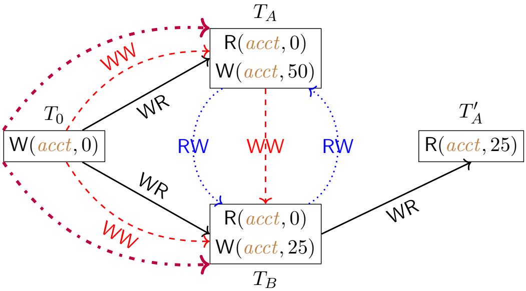

Dependency Graph-based Characterization of SI

Theorem (Theorem 4.1 of [Cerone and Gotsman, 2018])

For a history $\mathcal{H} = (T, \mathsf{SO})$,

$$\mathcal{H} \models \mathrm{SI} \iff \mathcal{H} \models \textsc{Int} \wedge$$
$$\exists\, \mathsf{WR}, \mathsf{WW}, \mathsf{RW}.\ \mathcal{G} = (\mathcal{H}, \mathsf{WR}, \mathsf{WW}, \mathsf{RW}) \wedge$$
$$(((\mathsf{SO}_{\mathcal{G}} \cup \mathsf{WR}_{\mathcal{G}} \cup \mathsf{WW}_{\mathcal{G}}) \ ; \ \mathsf{RW}_{\mathcal{G}}?) \ \textit{ is acyclic}).$$

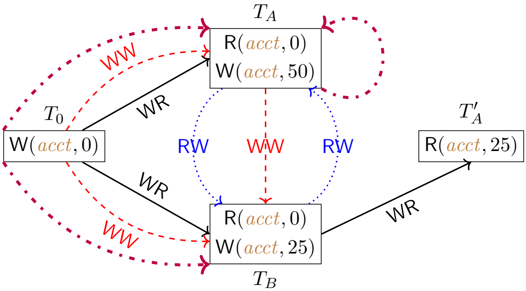

Dependency Graph-based Characterization of SI

Theorem (Theorem 4.1 of [Cerone and Gotsman, 2018])

For a history $\mathcal{H} = (T, \mathsf{SO})$,

$$\mathcal{H} \models \mathrm{SI} \iff \mathcal{H} \models \textsc{Int} \wedge$$
$$\exists\, \mathsf{WR}, \mathsf{WW}, \mathsf{RW}.\ \mathcal{G} = (\mathcal{H}, \mathsf{WR}, \mathsf{WW}, \mathsf{RW}) \wedge$$
$$(((\mathsf{SO}_{\mathcal{G}} \cup \mathsf{WR}_{\mathcal{G}} \cup \mathsf{WW}_{\mathcal{G}})\ ;\ \mathsf{RW}_{\mathcal{G}}?)\ \ \textit{is acyclic}).$$

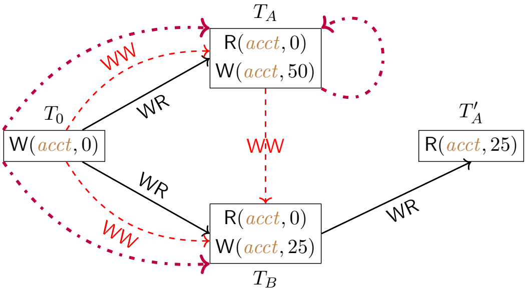

Dependency Graph-based Characterization of SI

$\mathcal{Q}$: How to capture all possible WW dependencies?

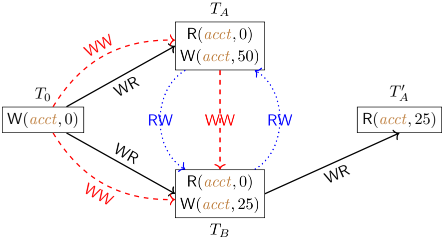

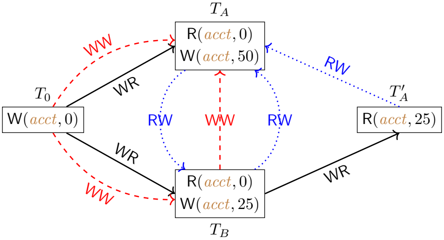

Dependency Graph-based Characterization of SI

$\mathcal{Q}$: How to capture all possible WW dependencies?

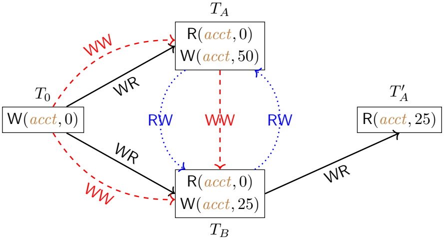

$\mathcal{A}$: encode them into SAT formulas based on (generalized) polygraphs

Generalized Polygraphs: A Family of Dependency Graphs

$$\begin{array}{ccc} T & & T' \\ \boxed{\mathsf{W}(x, 1)} & \xrightarrow{\mathsf{WR}(x)} & \boxed{\mathsf{R}(x, 1)} \\ & & \\ \boxed{\mathsf{W}(x, 2)} & \xrightarrow[\mathsf{WR}(x)]{} & \boxed{\mathsf{R}(x, 2)} \\ S & & S' \end{array}$$

$$\begin{array}{ccc} T & & T' \\ \boxed{\mathsf{W}(x, 1)} & \xrightarrow{\mathsf{WR}(x)} & \boxed{\mathsf{R}(x, 1)} \\ & & \\ \boxed{\mathsf{W}(x, 2)} & \xrightarrow[\mathsf{WR}(x)]{} & \boxed{\mathsf{R}(x, 2)} \\ S & & S' \end{array}$$

Generalized Polygraphs: A Family of Dependency Graphs

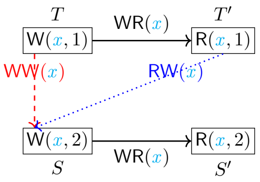

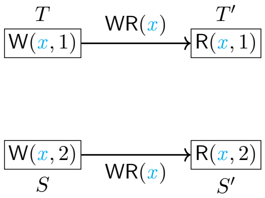

Generalized Polygraphs: A Family of Dependency Graphs

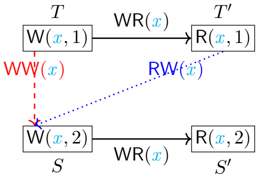

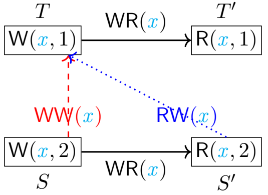

Generalized Polygraphs: A Family of Dependency Graphs

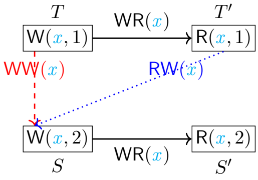

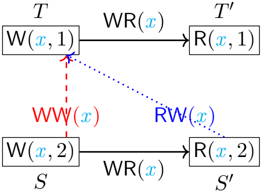

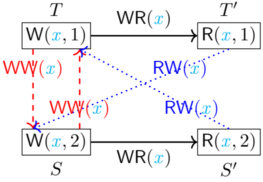

generalized polygraph:

$$\langle \mathit{either} \triangleq \{T \xrightarrow{\mathsf{WW}} S, T' \xrightarrow{\mathsf{RW}} S\}, \mathit{or} \triangleq \{S \xrightarrow{\mathsf{WW}} T, S' \xrightarrow{\mathsf{RW}} T\} \rangle$$

PolySI: An Illustrating Example of "Long Fork"

T_0 $\boxed{\mathsf{W}(x, 0)\ \mathsf{W}(y, 0)}$

PolySI: An Illustrating Example of "Long Fork"

T_1

$\boxed{\mathsf{W}(x, 1)}$

T_0 $\boxed{\mathsf{W}(x, 0)\ \mathsf{W}(y, 0)}$

PolySI: An Illustrating Example of "Long Fork"

T_1

$\boxed{\mathsf{W}(x, 1)}$

T_0 $\boxed{\mathsf{W}(x, 0)\ \mathsf{W}(y, 0)}$

$\boxed{\mathsf{W}(y, 1)}$

T_2

PolySI: An Illustrating Example of "Long Fork"

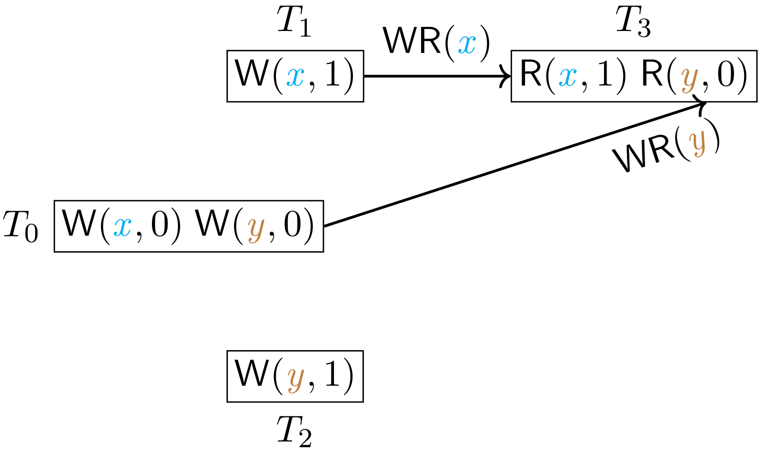

PolySI: An Illustrating Example of "Long Fork"

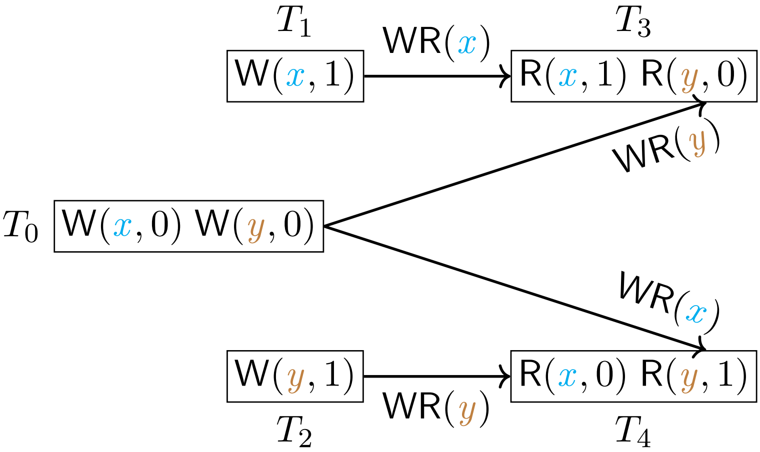

PolySI: An Illustrating Example of "Long Fork"

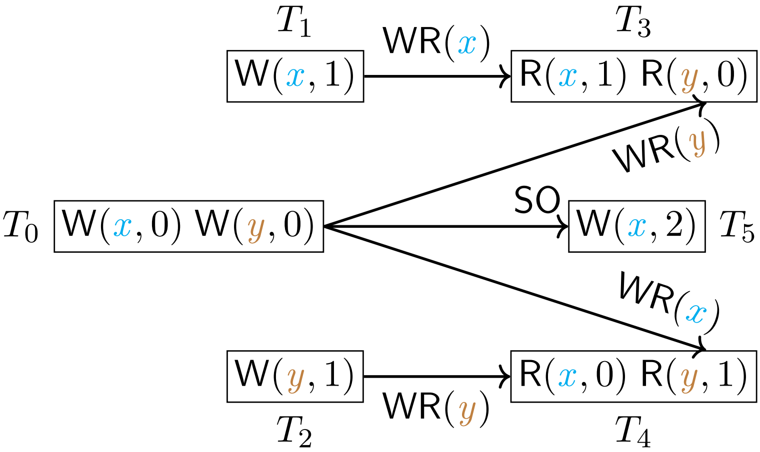

order between T_0, T_1, and T_5 (on x) and between T_0 and T_2 (on y)

PolySI: An Illustrating Example of “Long Fork”

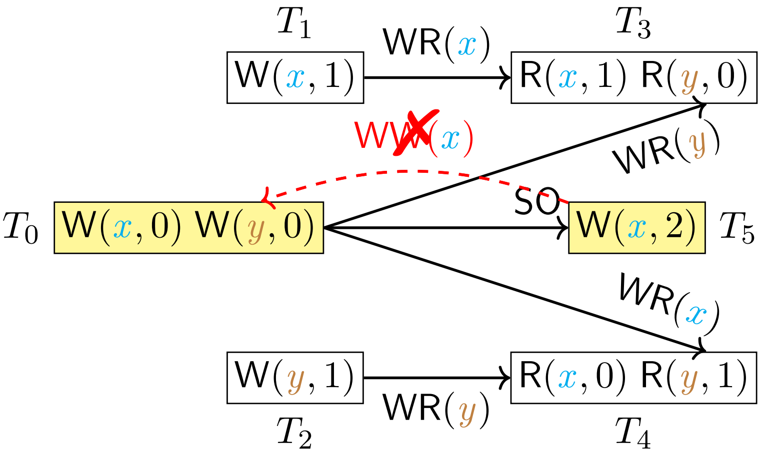

The $T_5 \xrightarrow{\text{WW}(x)} T_0$ case is pruned due to $T_0 \xrightarrow{\text{SO}} T_5 \xrightarrow{\text{WW}(x)} T_0$.

PolySI: An Illustrating Example of “Long Fork”

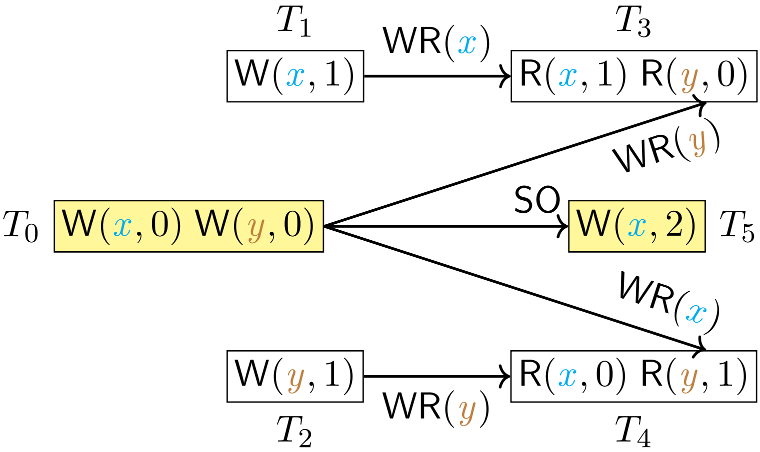

PolySI: An Illustrating Example of "Long Fork"

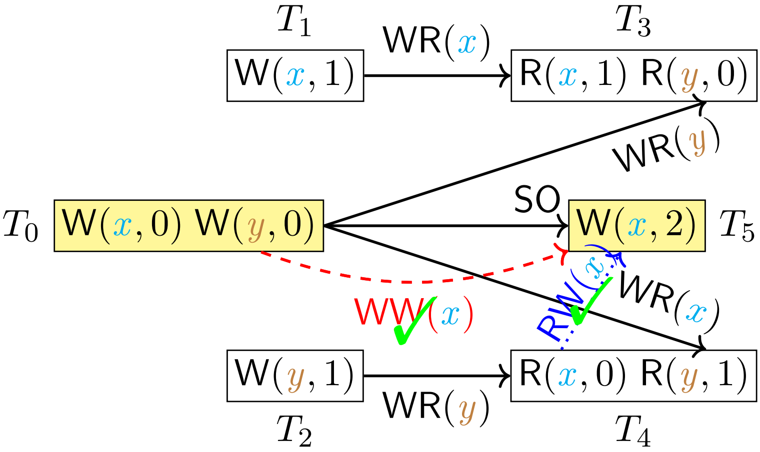

The $T_0 \xrightarrow{\mathsf{WW}(x)} T_5$ case becomes known.

PolySI: An Illustrating Example of "Long Fork"

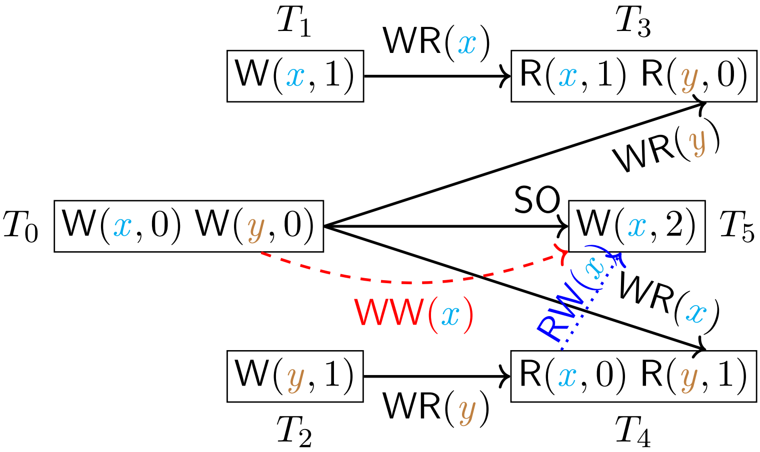

PolySI: An Illustrating Example of “Long Fork”

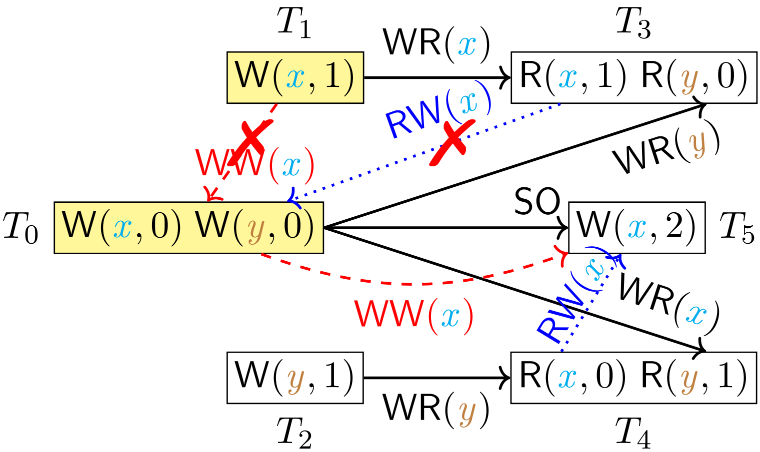

The $T_1 \xrightarrow{\mathsf{WW}(x)} T_0$ case is pruned due to $T_3 \xrightarrow{\mathsf{RW}(x)} T_0 \xrightarrow{\mathsf{WR}(y)} T_3$.

PolySI: An Illustrating Example of "Long Fork"

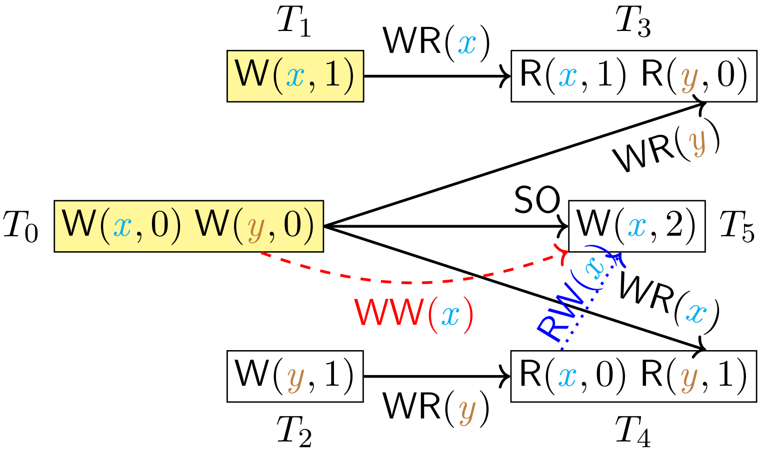

PolySI: An Illustrating Example of "Long Fork"

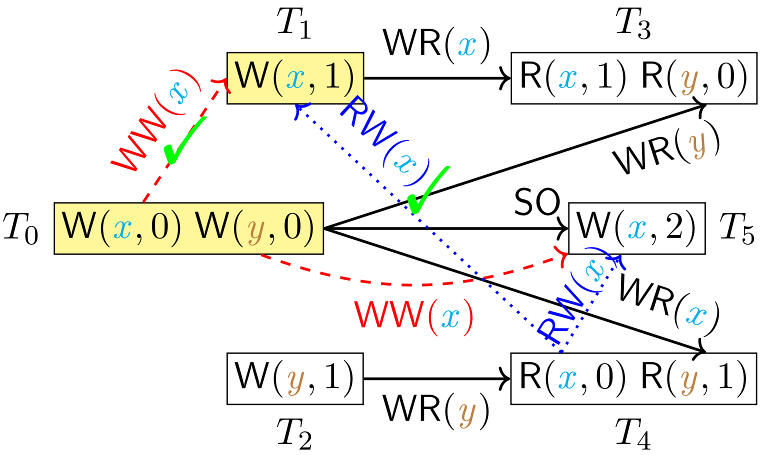

The $T_0 \xrightarrow{\mathsf{WW}(x)} T_1$ case becomes known.

PolySI: An Illustrating Example of “Long Fork”

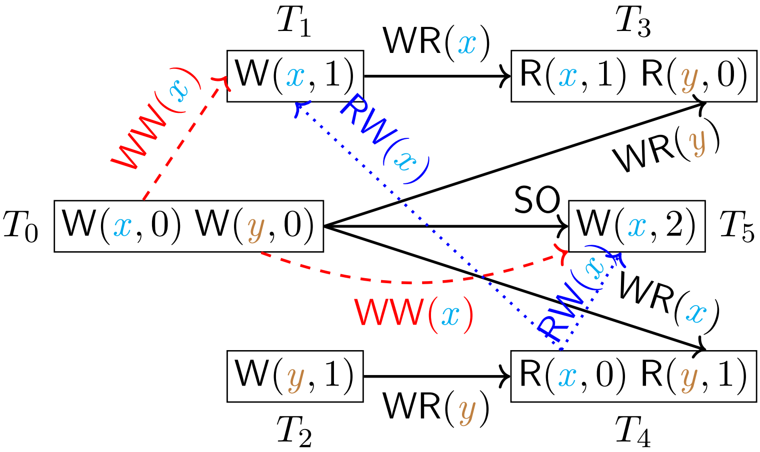

PolySI: An Illustrating Example of "Long Fork"

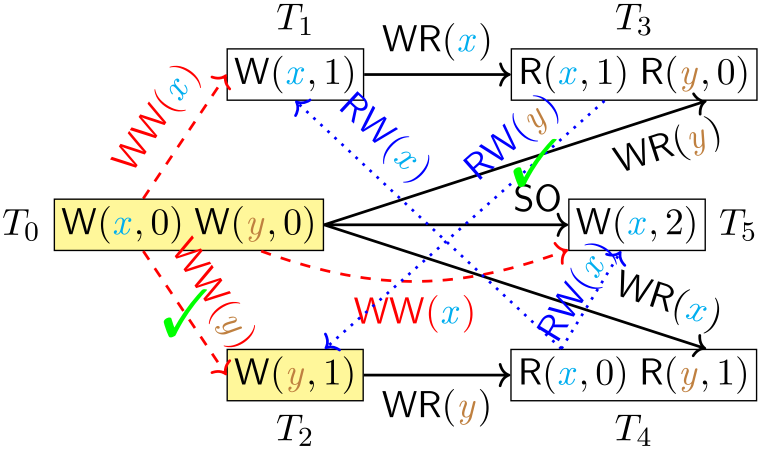

The $T_2 \xrightarrow{\mathsf{WW}(y)} T_0$ case is pruned,

while the $T_0 \xrightarrow{\mathsf{WW}(y)} T_2$ case becomes known.

PolySI: An Illustrating Example of "Long Fork"

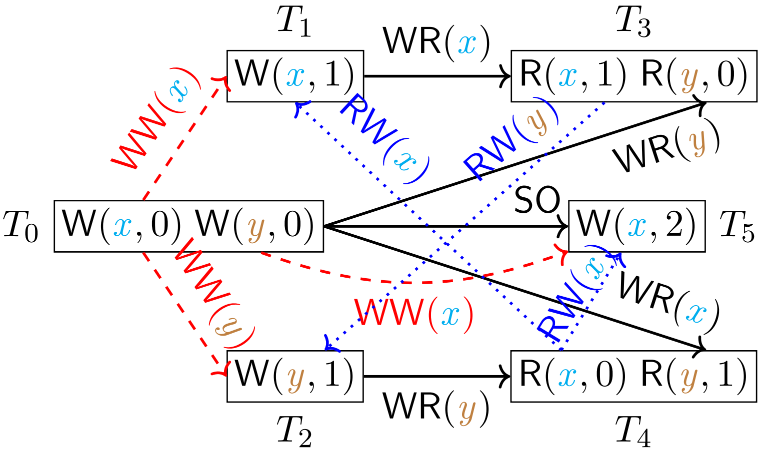

PolySI: An Illustrating Example of "Long Fork"

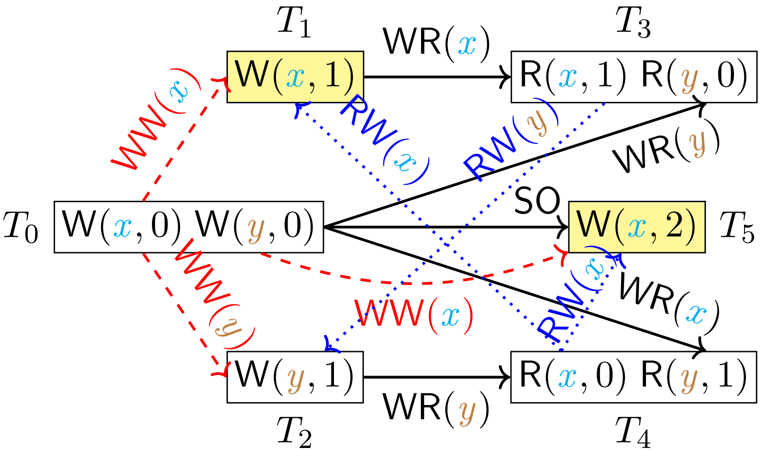

The order between T_1 and T_5 is still uncertain after pruning.

PolySI: An Illustrating Example of "Long Fork"

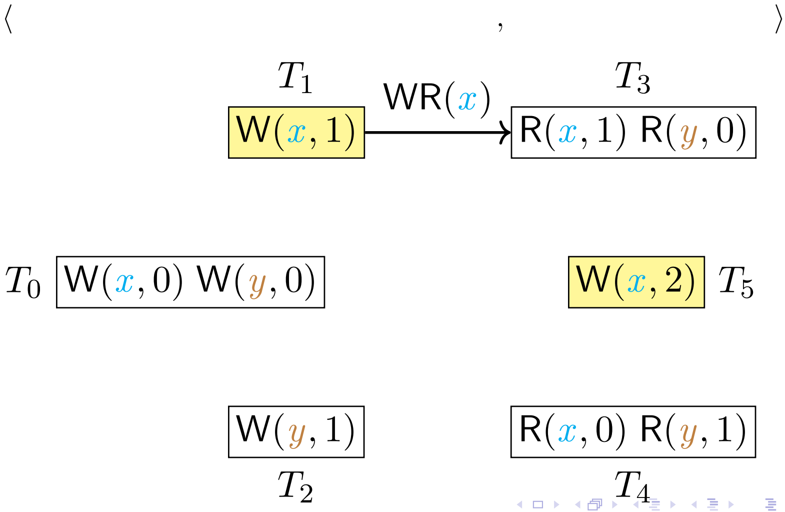

PolySI: An Illustrating Example of "Long Fork"

$\langle \mathit{either} = \{T_1 \xrightarrow{\mathsf{WW}(x)} T_5, T_3 \xrightarrow{\mathsf{RW}(x)} T_5\}, \qquad \rangle$

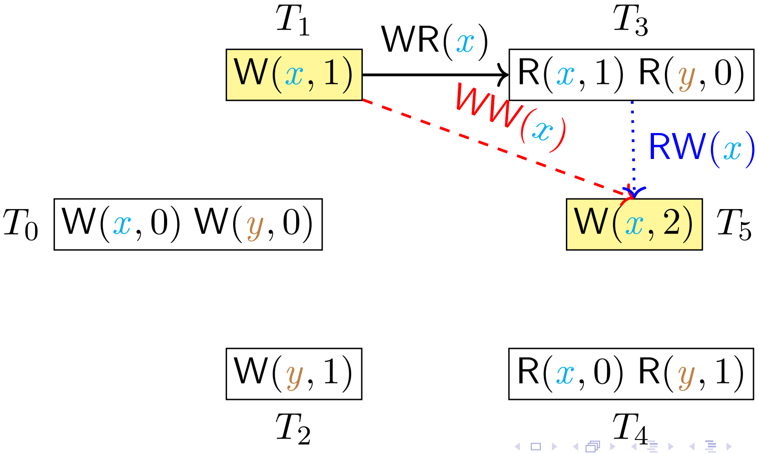

PolySI: An Illustrating Example of "Long Fork"

$$\langle \mathit{either} = \{T_1 \xrightarrow{\mathsf{WW}(x)} T_5, T_3 \xrightarrow{\mathsf{RW}(x)} T_5\}, \mathit{or} = \{T_5 \xrightarrow{\mathsf{WW}(x)} T_1\}\rangle$$

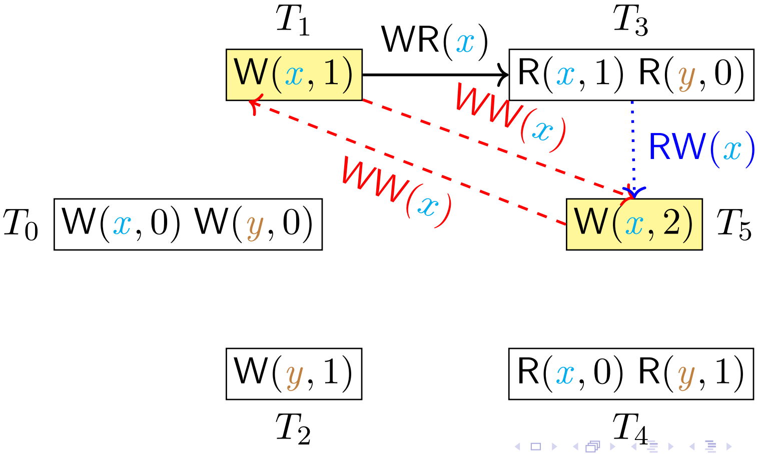

PolySI: An Illustrating Example of "Long Fork"

$$\langle either = \{T_1 \xrightarrow{\mathsf{WW}(x)} T_5, T_3 \xrightarrow{\mathsf{RW}(x)} T_5\}, or = \{T_5 \xrightarrow{\mathsf{WW}(x)} T_1\}\rangle$$

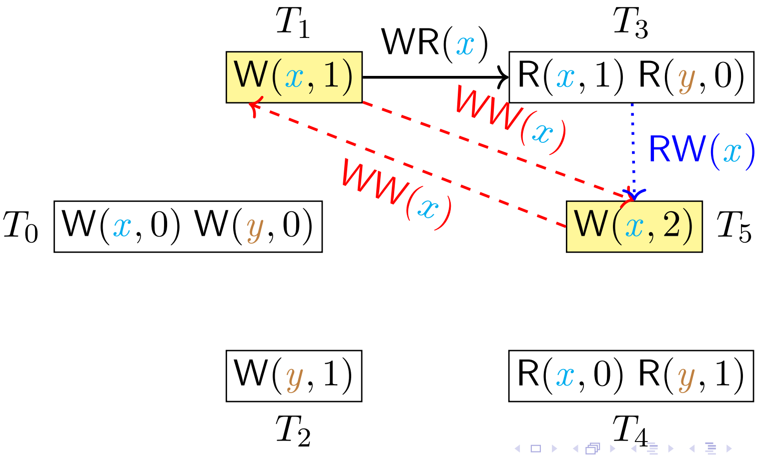

PolySI: An Illustrating Example of “Long Fork”

PolySI: An Illustrating Example of "Long Fork"

PolySI: An Illustrating Example of "Long Fork"

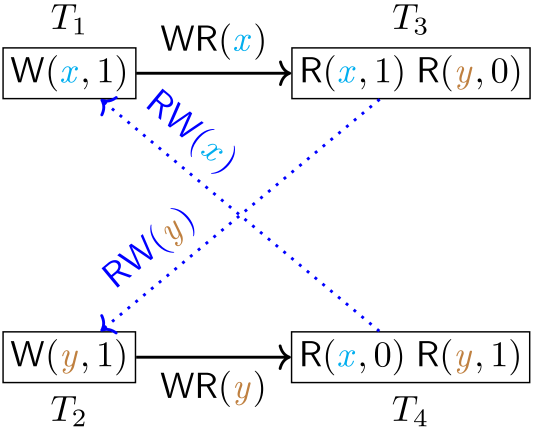

The undesired cycle for "long fork" found by MonoSAT.

Experimental Evaluation

(1) *Effective:* Can PolySI find SI violations in production databases?

(2) *Informative:* Can PolySI provide understandable counterexamples for SI violations?

(3) *Efficient:* How efficient is PolySI? Is it scalable?

Workloads, Benchmarks, and Setup

Finding SI Violations

Understanding Violations

Performance

Scalability

Conclusion

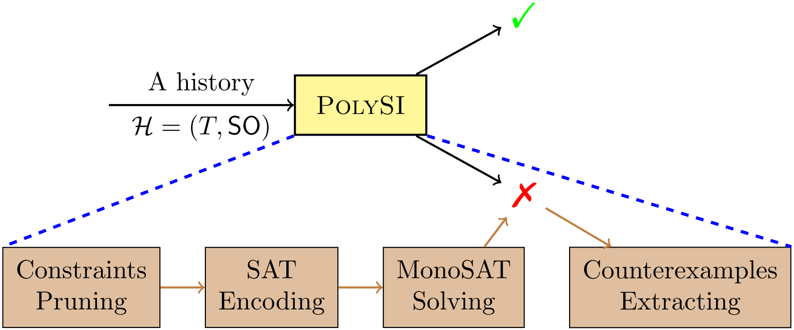

A history
$\mathcal{H} = (T, \mathsf{SO})$
PolySI
✓
✗
Constraints Pruning
SAT Encoding
MonoSAT Solving
Counterexamples Extracting

Hengfeng Wei (hfwei@nju.edu.cn)

Cerone, Andrea and Alexey Gotsman (Jan. 2018). “Analysing Snapshot Isolation”. In: *J. ACM* 65.2. ISSN: 0004-5411. DOI: 10.1145/3152396. URL: https://doi.org/10.1145/3152396.